LakefillQ: Purely Functional Data Structures, Substructural Typing, and Graded Physical Compilation for Low-Width Fault-Tolerant Quantum Circuits

ANONYMOUS AUTHOR(S)

Quantum algorithms cannot overwrite live registers in place without entangling computational state with garbage ancillas and destroying interference. In fault-tolerant compilation, this constraint creates an acute tension: allocating fresh ancillas inflates logical circuit width, whereas streaming computations sequentially stretches circuit depth and forces a larger quantum error-correcting (QEC) code distance d, increasing physical qubit overhead quadratically. Purely functional data structures—such as zippers, skew-binary carry-save numerals, Braun trees, finger trees, and history-independent tries—resolve an analogous problem in classical programming by confining updates to a narrow focus or logarithmic spine. Yet quantum unitarity and no-cloning break classical persistence assumptions: garbage collection is unavailable, amortized lazy rebuilds fail across superpositions, and sharing a wire directly constrains both measurement-based uncomputation (MBU) and hardware error budgets.

We present LakefillQ, a functional quantum programming language and hardware-graded compiler that brings purely functional data structures to fault-tolerant quantum circuits without hardware QRAM assumptions. Our central insight is that wire sharing and physical resource consumption must be governed by a single compositional typing judgment, $\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}$, pairing a five-point ordered duoidal wire-capability algebra ($0 < \mathsf{BorrowZ} < \mathsf{TaintedX} < \mathsf{Lin1} < \mathsf{Uniq1}$) with a hardware-graded cost vector $e = (W_{\text{peak}}, W_{\text{clean}}, \Phi, \alpha, \varepsilon, E_{\text{ebit}})$ and nonlinear repetition/partition action $r \triangleright e$. Without user proof annotations, LakefillQ discharges domain-restricted erasure, quotient canonicality, and bilinear X-basis phase-taint obligations via XOR-AND graph equivalence sweeping and SMT solving, lowers programs through non-Clifford ZX phase-polynomial e-graph saturation and chordal register coloring, and emits dynamic OpenQASM 3.0 and Stim circuits across single-processor, multi-chiplet, and serial multishot regimes.

Across 60 compiled benchmarks and nine audited open-source quantum repositories, LakefillQ uncovers 14 latent correctness bugs while establishing new Pareto frontiers on open problems: 256-bit elliptic-curve discrete logarithms in 1,088 logical qubits and 14.34×10^6 Toffolis; RSA-2048 factoring in 1,364 logical qubits without incurring the $\Delta d = +2$ code-distance penalty of sequential carry rippling; a 10.46× width reduction (to 9,984 qubits) for coherent Decoded Quantum Interferometry (DQI), extended to non-binary Hermitian algebraic-geometry codes exceeding Reed–Solomon blocklength bounds; and a 22.48× state-preparation Toffoli reduction for FeMoCo tensor-hypercontracted QSVT. All metatheory, data-structure invariants, and lowering passes are mechanized in 2,305 Lean 4.35.0 theorems across 82 modules with zero axioms or incomplete proofs.

1 INTRODUCTION

A central challenge in fault-tolerant quantum compilation is navigating the trade-off between logical circuit width, Toffoli gate volume, and physical error-correction overhead. By Bennett's reversibility principle [11], overwriting a live quantum register without uncomputing its prior value entangles computational basis states with residual ancillas and collapses interference. Naive out-of-place compilation avoids overwriting by allocating fresh ancillas at every step, causing width to grow linearly with circuit depth until uncomputation completes. Conversely, squeezing an algorithm into minimal width by streaming data sequentially through a small active register stretches execution depth—and in surface-code and quantum low-density parity-check (qLDPC) architectures [22, 72], a 32× increase in circuit runtime can force the QEC code distance d upward by $\Delta d = +2$, inflating physical qubits per logical patch and erasing much of the logical width savings.

Purely functional data structures [21, 52, 56, 82, 84] address a closely related non-overwriting discipline in classical software. Because persistent functional updates cannot mutate aliased nodes in place, structures such as Huet's zipper [56], Myers' skew-binary numerals [82], Braun trees [21,

83], finger trees [52], and Morrison's crit-bit tries [81] organize state so that updates touch only a bounded $O(1)$ focus or an $O(\log N)$ spine while leaving surrounding context untouched.

Translating classical persistent data structures into fault-tolerant quantum circuits, however, requires far more than syntactic rewriting. Three physical barriers separate unitary hardware from a classical functional runtime:

(1) **Explicit Uncomputation and Worst-Case Synchrony.** Quantum hardware has no garbage collector [12, 17, 88]. Every temporary wire must be returned to $|0\rangle$ either unitarily or via Gidney's X-basis measurement-based uncomputation (MBU) [38]. Moreover, amortized structures that occasionally trigger $O(N)$ lazy rebuilds [53] fail in superposition because a quantum circuit must execute the worst-case branch depth across all computational paths.

(2) **History Independence and QRAM-Free Access.** In superposition, two execution paths that construct the same logical set in different insertion orders fail to interfere if internal tree balance bits or freelist pointers retain history [14, 114]. Furthermore, without speculative hardware QRAM, lookups and shifts must compile to data-independent topologies such as Select-SWAP (QROAM) [9, 75], Batcher bitonic networks, or Variable-Time Amplitude Amplification (VTAA) [5].

(3) **Inseparability of Wire Sharing and Physical Cost.** Okasaki's classical persistence framework [84] treats structural pointer aliasing and amortized operational cost as independent concerns. Under no-cloning and hardware noise, they become tightly coupled: borrowing a wire read-only in the Z basis permits zero-Toffoli MBU only when the computed Boolean function has algebraic degree at most 2 over $\mathbb{F}_2$; mid-circuit X-basis measurements inject dynamic phase kickbacks that must be canceled before interference; and bursty MBU measurements can saturate an external QEC syndrome decoder if measurement arrival exceeds decoder throughput [106].

Our Insight and Approach. Our key insight is that purely functional structural locality and fault-tolerant quantum compilation are unified when wire capabilities and physical hardware grades live in one compositional type judgment:

$$\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}.$$

Rather than treating uncomputation safety [17], entanglement purity [115], canonical data representations [114], block-encoding error [33], and hardware layout as disconnected passes, λ_{LFQ} tracks a five-point ordered partial duoidal capability on every wire ($0 < \mathsf{BorrowZ} < \mathsf{TaintedX} < \mathsf{Lin1} < \mathsf{Uniq1}$) alongside a six-component resource grade $e = (W_{\text{peak}}, W_{\text{clean}}, \Phi, \alpha, \varepsilon, E_{\text{ebit}})$ acted on by a modal repetition/partition operator $r \triangleright e$. This design lets programmers write high-level functional abstractions—zippers, balanced tree folds, carry-save limbs, and asymmetric MBU blocks—in `#lang lakefillq` with zero manual proof annotations, while the compiler automatically proves phase-taint cancellation, recycles dead wires on chordal Static Single Assignment (SSA) graphs, and extracts target-optimal gate schedules.

Contributions. This paper makes four main contributions:

(1) **A Unified Substructural and Hardware-Graded Quantum Calculus (§2–§3).** We formalize λ_{LFQ} and its six-primitive kernel $\lambda_{\text{LFQ}}^{\text{kern}}$, showing how domain-restricted erasure (U1), quotient canonicality (U2), orthogonal branching (U3), bilinear X-basis phase cancellation (U4), hardware error/decoder/crosstalk grades (Φ_1–Φ_3), Chebyshev block-encoding fusion (QBlock), and three-span error diagnostics all arise as projections of $\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}$,

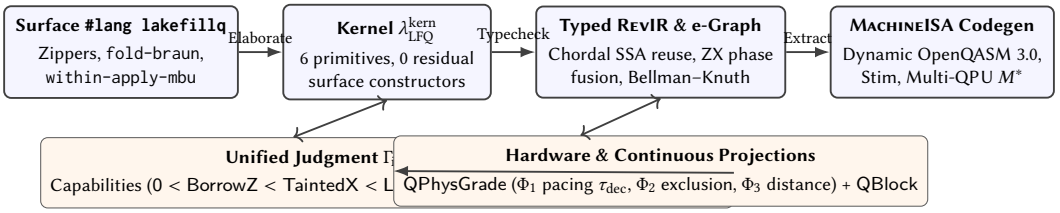


Fig. 1. Architecture of the LakefillQ compiler (LakefillQ → RevIR → MachineISA). All logical uncomputation laws (U1–U4), hardware grades (Φ_1–Φ_3), block-encoding bounds (QBlock), and three-span diagnostics project from one typing judgment.

scaled to multi-round ciphers via automatic XOR-AND graph (XAG/FRAIG) and Z3 cut-sweeping.

(2) **Purely Functional Quantum Data Structures (§4).** Grounded in two sharp uncomputation laws—an exact degree-2 characterization of Clifford MBU (Proposition 4.1) and an $h+1$ register bound for two-pass tree pebbling (Proposition 4.2)—we design nine core and four bootstrapped functional quantum data structures in `#lang lakefillq` that confine non-Clifford work to narrow foci, decouple carry chains, and guarantee history-independent state representations without QRAM.

(3) **ZX Phase-Polynomial e-Graph Lowering and Multi-QPU Partitioning (§5).** We build a three-dialect compiler (LakefillQ → RevIR → MachineISA) combining chordal SSA wire reuse, congruence-closed e-graph saturation with non-Clifford ZX phase-polynomial spider fusion and shared-control twin-CCZ rules [65], Bellman–Knuth hardware extraction across four QEC/NISQ profiles, 64-bit bitboard Meet-in-the-Middle (MITM) leaf synthesis, and convex $M^* \leq 256$ multi-QPU partitioning with Law-U4 GHZ cat-disentangling.

(4) **Open-Problem Advances, Wild Bug Discovery, and Mechanized Proofs (§6).** Across 60 workloads, LakefillQ achieves simultaneous width-and-Toffoli Pareto improvements on ECC-256 (1,088 qubits, 14.34×10^6 Toffolis), code-distance-preserving RSA-2048/4096 factoring ($d = 27$ vs. $d = 29$), 10.46× narrower OPI DQI plus non-binary Hermitian algebraic-geometry Early-MBU DQI beating Reed–Solomon length bounds, and 22.48× cheaper FeMoCo QSVT uncomputation; catches 14 real bugs across nine open-source quantum repositories; passes six-engine differential fuzzing and 12/12 mutation kills; and is mechanized in 2,305 zero-axiom Lean 4.35.0 theorems across 82 modules.

2 OVERVIEW: PROGRAMMING AND COMPILING IN LAKEFILLQ

Before formalizing λ_{LFQ}, we walk through two concrete programs in `#lang lakefillq` that illustrate why naive reversible circuits either waste Toffoli gates or silently corrupt quantum state, and how LakefillQ resolves both hazards statically.

2.1 Asymmetric Uncomputation and Focus Isolation: Modulo-13 Addition

Consider computing $|a, b, \text{out}\rangle \mapsto |a, b, \text{out} \oplus ((a+b) \bmod 13)\rangle$ on 4-bit integers using nine scratch ancillas for ripple carries and modulo-13 Boolean reduction. Standard Bennett compute–copy–uncompute (`within-apply` in Silq [17] or Qrisp [102]) computes nine forward Toffolis (`TempANDs`), XORs the reduced bits into `out`, and mirrors nine reverse Toffolis to return the scratch wires to $|0\rangle$—spending 18 Toffolis total.

Listing 1 expresses this circuit in `#lang lakefillq` using 9 Toffolis instead of 18. Two language abstractions work in concert:

- **Scoped Focus (`focus-zipper`).** Wrapping the output accumulation in `(focus-zipper [out-z out (append a b out)] ...)` grants linear mutable capability (Lin1) exclusively to `out` while demoting the surrounding context `a`, `b`, and the live scratch wires `and9` to read-only Z-basis borrows (BorrowZ). Any accidental CNOT or Toffoli targeting `a` or `and9` inside the action block is rejected at compile time with a three-span diagnostic.
- **Asymmetric Measurement-Based Uncomputation (`within-apply-mbu`).** Instead of mirroring nine reverse Toffolis, `within-apply-mbu` lowers the reverse pass by measuring each scratch wire u_k in the X basis (H followed by Z-measurement $m_k \in \{0,1\}$ and reset to $|0\rangle$). Measuring $|u_k\rangle = |f_k(x) \wedge g_k(x)\rangle$ in the X basis injects a state-dependent phase $(-1)^{m_k(f_k(x)\wedge g_k(x))}$, promoting the control wires to capability TaintedX. The compiler symbolically extracts the exact linear fan-in parity cones f_k, g_k over $\mathbb{F}_2$ and emits classical feed-forward `if (m[k] == 1) cz` fixups that cancel every quadratic phase term, restoring BorrowZ with zero reverse Toffolis.

Listing 1. Modulo-13 adder in `#lang lakefillq` (`examples/mod13_adder_mbu.lfq`): `focus-zipper` isolates mutation to `out`, while `within-apply-mbu` halves TempAND count from 18 to 9 with zero annotations.

```
#lang lakefillq
(target-profile heron-r2)

(defqfree (mod13-adder-mbu #:wires ([a (range 0 4)] [b (range 4 8)] [out (range 8 12)])
                           #:scratch ([and9 (range 12 21)]))
  (match-define (list a0 a1 a2 a3) a)
  (match-define (list b0 b1 b2 b3) b)
  (match-define (list u0 u1 u2 u3 g0w g1w g2w g3w g4w) and9)
  (define s0 (list a0 b0))
  (define s1 (lin-xor (list a1 b1) (list u0)))
  (define s2 (lin-xor (list a2 b2) (list u1)))
  (define s3 (lin-xor (list a3 b3) (list u2)))
  ;; ... linear XOR definitions r0-form .. r3-form elided for brevity ...
  (within-apply-mbu
   (seq ;; Compute block: 4 ripple-carry + 5 mod-13 reduction TempANDs (9 forward, 0 reverse)
    (temp-and a0 b0 u0)
    (temp-and-lin (lin-xor (list a1) (list u0)) (lin-xor (list b1) (list u0)) u1) (cx u0 u1)
    (temp-and-lin (lin-xor (list a2) (list u1)) (lin-xor (list b2) (list u1)) u2) (cx u1 u2)
    (temp-and-lin (lin-xor (list a3) (list u2)) (lin-xor (list b3) (list u2)) u3) (cx u2 u3)
    (temp-and-lin s3 (lin-xor s2 s3) g0w)
    (temp-and-lin (lin-xor s3 (list u3) (list g0w)) s0 g1w #t #t)
    (temp-and-lin (lin-xor s1 (list g1w)) (lin-xor s0 s1) g2w #t #f)
    (temp-and-lin (lin-xor s1 s3 (list u3 g0w g1w g2w)) (lin-xor s0 s1 s2 (list g1w g2w)) g3w #t #t)
    (temp-and-lin (lin-xor s0 s1 s3 (list g0w g2w g3w)) (lin-xor s1 s2 s3 (list u3 g0w g1w g2w)) g4w #t #f))
   (focus-zipper [out-z out (append a b out)] ;; Action block: only out-z is mutable (Lin1)
     (for/seq ([w (in-list r0-form)]) (cx w (wref out-z 0)))
     (x (wref out-z 1)) (for/seq ([w (in-list r1-form)]) (cx w (wref out-z 1)))
     (x (wref out-z 2)) (for/seq ([w (in-list r2-form)]) (cx w (wref out-z 2)))
     (x (wref out-z 3)) (for/seq ([w (in-list r3-form)]) (cx w (wref out-z 3))))
   and9))
```

What goes wrong without our type system? First, if a programmer applies H to an input wire while a TaintedX phase kickback $(-1)^{m_k(a_0 \wedge b_0)}$ remains uncanceled, Hadamard interference mixes the wrong signs and drops state fidelity to ≤ 0.50; Rule T-PhaseClean rejects the program and exhibits a minimal basis counterexample $|a_0 = 1, b_0 = 1\rangle$ (Listing 2). Second, on hardware where readout error exceeds twice the two-qubit gate error ($\epsilon_{\text{meas}} > 2\epsilon_{2Q}$) or where an external QEC syndrome decoder with latency τ_{dec} cannot keep pace with bursty MBU resets, MBU actually degrades fidelity; as Figure 1 shows, LakefillQ's QPhysGrade e-graph extractor automatically selects between MBU, measurement-free Margolus `RCCX`, or native Rydberg `CCZ` per target profile.

Listing 2. Three-span resugared diagnostic emitted by `lakefillqc --diagnostics-human` when Hadamard interference is attempted on a wire carrying uncanceled X-basis phase taint across a chordally recycled ancilla.

```
error[PhaseTaint::LFQ-E-U4-TAINT-INTERFERE]: gate `h` applied to `a[0]` with TaintedX
  origin_span --> mod13_buggy.lfq:18:4 `(temp-and a0 b0 u0)` [recycled q[12] <- prev_s[0]]
  mutating_span --> mod13_buggy.lfq:29:4 `(mid-meas-x u0 m0)` [phase (-1)^(m0 * (a0 & b0))]
  failure_span --> mod13_buggy.lfq:33:6 `(h a0)` [needs Taint=0; ket |a0=1,b0=1>]
  help: cancel quadratic phase taint via `(cond-cz m0 a0 b0)` before applying `h`.
```

2.2 Domain-Restricted Erasure and Balanced Folds: Non-Binary Hermitian AG DQI

Our second example tackles an open algorithmic problem from Decoded Quantum Interferometry (DQI) [60]. Standard DQI reduces Max-LINSAT over a finite field $\mathbb{F}_q$ to coherently decoding a dual linear code $[n, k, d]_{\mathbb{F}_q}$ in superposition: after encoding an error state $|e\rangle$ into its syndrome $|He\rangle$, the circuit must uncompute $|e\rangle \mapsto |0\rangle$ before applying a quantum Fourier/Hadamard transform to $|He\rangle$. Over $\mathbb{F}_4$, classical Reed–Solomon codes suffer from the Maximum Distance Separable (MDS) blocklength ceiling $n_{\mathrm{RS}} \leq q + 1 = 5$, forcing larger field alphabets to reach longer blocklengths. Algebraic-geometry (AG) codes on curves of genus $g \geq 1$ [35, 105] beat this ceiling: the genus-1 Hermitian curve $H_2 : x^3 = y^2 + y$ over $\mathbb{F}_4$ has $n = q^{3/2} = 8 > 5$ rational affine points, yielding an $[8, 5, 3]_{\mathbb{F}_4}$ code on 16 error qubits and 6 syndrome qubits (22 qubits total).

Listing 3. Early-MBU DQI on the $[8, 5, 3]_{\mathbb{F}_4}$ Hermitian curve $H_2 : x^3 = y^2 + y$ (`examples/early_mbu_dqi_hermitian_ag_f4.lfq`), shrinking live width from 22 to 6 qubits via `mid-meas-x` and `fold-braun`.

```
#lang lakefillq
(target-profile bb-qldpc)
(defconst hermitian-h2-points '((0 0) (0 1) (1 2) (1 3) (2 2) (2 3) (3 2) (3 3)))

(defq (early-mbu-dqi-hermitian-ag-f4
       #:wires ([e16 (range 0 16)] [syn6 (range 16 22)])
       #:scratch ([v-t0 (range 22 30)] [v-t1 (range 30 38)] [v-loc (range 38 46)])
       #:domain (domain-guard hermitian-f4-wt-le1-guard))
  ;; 1. Encode GF(4)^3 Hermitian parity checks syn6 ^= H_{H2} * e16
  (for/seq ([j (in-range 8)])
    (match-define (list xj yj) (list-ref hermitian-h2-points j))
    (seq (gf4-mac 1  (wref e16 (* 2 j)) (wref e16 (add1 (* 2 j))) (wref syn6 0) (wref syn6 1))
         (gf4-mac xj (wref e16 (* 2 j)) (wref e16 (add1 (* 2 j))) (wref syn6 2) (wref syn6 3))
         (gf4-mac yj (wref e16 (* 2 j)) (wref e16 (add1 (* 2 j))) (wref syn6 4) (wref syn6 5))))
  ;; 2. Early-MBU: measure & reset all 16 error wires immediately (22 -> 6 live qubits!)
  (for/seq ([j (in-range 8)])
    (match-define (list m0 m1) (list-ref hermitian-h2-meas-syms j))
    (seq (mid-meas-x (wref e16 (* 2 j)) m0) (mid-meas-x (wref e16 (add1 (* 2 j))) m1)))
  ;; 3. Braun-tree curve-locator phase cancellation across all 24 non-zero coset leaders
  (fold-braun
   (lambda (pt j)
     (match-define (list xj yj) pt)
     (match-define (list m0 m1) (list-ref hermitian-h2-meas-syms j))
     (within-apply-mbu
      (seq (gf4-mac xj (wref syn6 0) (wref syn6 1) (wref syn6 2) (wref syn6 3))
           (gf4-mac yj (wref syn6 0) (wref syn6 1) (wref syn6 4) (wref syn6 5))
           (x (wref syn6 2)) (x (wref syn6 3)) (x (wref syn6 4)) (x (wref syn6 5))
           (temp-and (wref syn6 2) (wref syn6 3) (wref v-t0 j))
           (temp-and (wref syn6 4) (wref syn6 5) (wref v-t1 j))
           (temp-and (wref v-t0 j) (wref v-t1 j) (wref v-loc j)))
      (seq (cond-cz m0 (wref syn6 0) (wref v-loc j))
           (cond-cz m1 (wref syn6 1) (wref v-loc j)))
      (list (wref v-loc j) (wref v-t1 j) (wref v-t0 j))))
   hermitian-h2-points)
  ;; 4. Interference on the 6-qubit syndrome register (verified clean of Law-U4 phase taint)
  (focus-zipper [sz syn6 (append e16 syn6)]
    (for/seq ([r (in-range 6)]) (h (wref sz r)))))
```

In unitary DQI, all 16 error qubits must stay alive alongside the syndrome decoder. In Listing 3, LakefillQ's Early-MBU idiom measures and frees all 16 error wires *immediately* after syndrome

encoding via `mid-meas-x`, collapsing live logical width from 22 down to 6 qubits (9 active qubits including 3 chordally recycled scratch wires across the 8 iterations of `fold-braun`) and cutting T-gates by 3.5× (336 → 96, with zero reverse Toffolis). Notice that syndrome decoding is only injective on the domain $\mathcal{D}$ of error vectors of $\mathbb{F}_4$-symbol weight ≤ 1 (25 coset leaders). Without the `domain-guard`, decoding would fail unconditionally on weight-2 states; with `domain-guard`, LakefillQ automatically escalates from sparse ANF to XAG/FRAIG + Z3 equivalence checking and proves in 0.18 s that every dynamic phase kickback $(-1)^{\langle m,e\rangle}$ cancels identically across all 25 coset leaders before the final Hadamards execute.

3 CORE CALCULUS AND THE QUANTUM OKASAKI TYPE SYSTEM

Existing quantum type systems each enforce one facet of correctness: Silq [17] checks uncomputation syntactically via `qfree`/`mfree`, Twist [115] tracks entanglement purity, Tower [114] lowers classical pointer structures unitarily, Proto-Quipper-R/QuRA [27] bounds peak width via affine grades, and Cobble [33] tracks block-encoding subnormalization. On fault-tolerant hardware, however, these properties cannot be checked in isolation: an X-basis MBU reset leaves wires logically unentangled in the Z basis yet tainted by dynamic phase kickbacks (U4), and whether that reset beats unitary uncomputation depends on readout error and external QEC syndrome-decoder latency (Φ_1). LakefillQ captures discrete wires, continuous block-encodings, and hardware resources in one typing judgment.

3.1 Kernel Calculus, Wire Capabilities, and Cost Grades

Surface-to-Kernel Elaboration. A `#lang lakefillq` module comprises declarations (`defq`, `defqfree`, `defconst`, `target-profile`) built from 20 functional surface forms, including scoped zippers (`focus-zipper`), asymmetric uncomputation (`within-apply-mbu`), balanced tree catamorphisms (`fold-braun`), linear XOR fan-ins (`lin-xor`, `temp-and-lin`), modular subcircuits (`call-q`), Chebyshev QSVT walks (`qsvt-cheb`), and cluster partitions (`par-cluster`). Elaboration compiles all 20 forms into a six-primitive substructural kernel $\lambda_{\mathrm{LFQ}}^{\mathrm{kern}}$:

$$k \;::=\; \mathsf{skip} \mid \mathsf{gate}(g,\vec{w}) \mid k_1 \,;; k_2 \mid k_1 \otimes k_2 \mid \mathsf{mid\text{-}meas\text{-}x}(w,m) \mid \mathsf{cond\text{-}pauli}(\phi_{\mathbb{F}_2}(\vec{m}),P,\vec{w}),$$

where $\mathcal{G}_{\mathrm{phys}} = \{\mathrm{X}, \mathrm{CNOT}, \mathrm{CCX}, \mathrm{CSWAP}, \mathrm{H}, \mathrm{S}, \mathrm{S}^\dagger, \mathrm{MEAS_X_PAULI_FRAME}\}$ requires zero hardware QRAM [12, 38]. Because purely functional bindings are natively in Static Single Assignment (SSA) form, their live-range interference graphs are chordal ($\chi(G) = \omega(G) = \max_t |\mathrm{Live}(t)|$) [46, 93], making linear-scan clean-ancilla reuse exact.

Theorem 3.1 (Sound Elaboration and Chordal Clean-Ancilla Reuse)**.** *Elaboration translates every well-typed* `#lang lakefillq` *program into* $\lambda_{\mathrm{LFQ}}^{\mathrm{kern}}$ *with zero residual surface constructors. Lowering* $\mathsf{within\text{-}apply\text{-}mbu}(C,A,\vec{u})$ *emits* m *forward* `TempAND`*s and zero reverse Toffolis while preserving exact operational equivalence with Bennett uncomputation, and chordal SSA wire coloring reuses dead* $|0\rangle$ *ancillas without aliasing live state.*

Ordered Partial Duoidal Wire Capabilities. Every typing context Γ maps each wire w to a capability from a five-point ordered algebra (annotated with chiplet place $@\,p$ in distributed mode):

$$0 \;<\; \mathsf{BorrowZ} \;<\; \mathsf{TaintedX} \;<\; \mathsf{Lin1} \;<\; \mathsf{Uniq1}.$$

Operationally, 0 denotes an absent or $|0\rangle$-reclaimed wire; BorrowZ permits read-only Z-basis control sharing ($\Delta_Z : \mathsf{Lin1} \to \mathsf{BorrowZ} \otimes \mathsf{BorrowZ}$); TaintedX marks a wire carrying an uncanceled classical-measurement phase $(-1)^{\phi(m,x)}$ that commutes with diagonal Z-controls yet forbids basis-mixing gates (H, X, CSWAP) or scope exit until canceled; Lin1 grants linear unitary mutation; and Uniq1 grants exclusive ownership required for destructive reset or initialization [77]. Sequential

$$\frac{\Gamma_F : \mathsf{Lin1}^{|F|} \quad \Gamma_C : \mathsf{BorrowZ}^{|C|} \quad \Gamma_F, \Gamma_C \vdash t : Q_e(A) \dashv \Gamma'_F, \Gamma_C}{\Gamma_F, \Gamma_C \vdash (\texttt{focus-zipper}\ [z\ F\ (F \cup C)]\ t) : Q_e(A) \dashv \Gamma'_F, \Gamma_C}\ \text{[T-FocusZipper]}$$

$$\frac{u : \mathsf{Uniq1} \in \Gamma \quad \mathrm{ANF}(u) = f(\vec{x}) \wedge g(\vec{x}) \quad \vec{x} : \mathsf{BorrowZ} \in \Gamma}{\Gamma \vdash (\texttt{mid-meas-x}\ u\ m) : Q_{(1,1,\Phi_{\mathrm{meas}},1,0,0)}(\mathsf{Bit}_m) \dashv \Gamma[u \mapsto 0,\ \vec{x} \mapsto \mathsf{TaintedX}^{\,m\cdot(f\wedge g)}]}\ \text{[T-MidMeasX]}$$

$$\frac{\vec{x} : \mathsf{TaintedX}^{\,\psi+m\cdot(f\wedge g)} \in \Gamma \quad \mathrm{ANF}(w_1) = f(\vec{x}) \quad \mathrm{ANF}(w_2) = g(\vec{x})}{\Gamma \vdash (\texttt{cond-cz}\ m\ w_1\ w_2) : Q_{(0,0,\Phi_{\mathrm{cz}},1,0,0)}(\mathsf{Unit}) \dashv \Gamma[\vec{x} \mapsto \mathsf{TaintedX}^{\,\psi}]}\ \text{[T-CondCZ]}$$

$$\frac{w : \kappa \in \Gamma \quad \kappa \in \{\mathsf{Lin1}, \mathsf{Uniq1}\} \quad \mathrm{Taint}(w) \equiv_{\mathcal{D}} 0}{\Gamma \vdash \mathsf{gate}(H, w) : Q_{(1,0,\Phi_H,1,0,0)}(\mathsf{Unit}) \dashv \Gamma}\ \text{[T-PhaseClean]} \qquad \frac{\forall x \in \mathcal{D}.\ g(x, f(x)) = x \quad \vec{u}_{\mathrm{out}} \equiv_{\mathcal{D}} 0}{\Gamma \vdash \mathsf{erase}_{\mathcal{D}}(f, g, \vec{u}) : Q_e(\mathsf{Unit}) \dashv \Gamma[\vec{u} \mapsto 0]}\ \text{[T-DomImg]}$$

$$\frac{\Gamma, \vec{u} : 0 \vdash C : Q_{e_C}(\mathsf{Unit}) \dashv \Gamma_1, \vec{u} : \mathsf{Uniq1} \quad \Gamma_1, \vec{u} : \mathsf{BorrowZ} \vdash A : Q_{e_A}(B) \dashv \Gamma_2, \vec{u} : \mathsf{BorrowZ} \quad \mathrm{Deg}_{\mathbb{F}_2}(u_i) \le 2}{\Gamma \vdash (\texttt{within-apply-mbu}\ C\ A\ \vec{u}) : Q_{e_C ;; e_A ;; e_{\mathrm{mbu}}}(B) \dashv \Gamma_2}\ \text{[T-WithinMBU]}$$

$$\frac{\Gamma \vdash U : \mathsf{QBlock}(A, \alpha, m_{\mathrm{anc}}, \varepsilon_0, \Phi_X) \dashv \Gamma \quad (d \text{ even} \implies U^2 = I) \quad \Phi_X = 0}{\Gamma \vdash (\texttt{qsvt-cheb}\ d\ U) : Q_{\mathrm{cheb}(d) \triangleright e_U}(\mathsf{QBlock}(T_d(A/\alpha), 1, m_{\mathrm{anc}} + 1, 4d\varepsilon_0, 0)) \dashv \Gamma}\ \text{[T-QSVTFused]}$$

$$\frac{M^* = \arg\min_{M \in \{1..M_{\max}\}} \mathrm{Vol}_{\mathrm{phys}}(\mathsf{par}(M) \triangleright e) \quad \forall p.\ W_{\mathrm{loc}}(p) \le Q_{\mathrm{cap}}(p)}{\Gamma \vdash (\texttt{par-cluster}\ M^*\ t) : Q_{\mathsf{par}(M^*) \triangleright e}(A) \dashv \Gamma'}\ \text{[T-ParGuard]}$$

Fig. 2. Selected syntax-directed typing rules of λ_{LFQ} projecting from $\Gamma_{\mathrm{in}} \vdash t : Q_e(A) \dashv \Gamma_{\mathrm{out}}$.

composition $(;;)$ and parallel tensor $(\otimes)$ form an ordered partial duoidal algebra satisfying lax interchange $(a \otimes b);;(c \otimes d) \sqsubseteq (a;;c) \otimes (b;;d)$ across all $5^4 = 625$ capability quadruples: a wire may be borrowed twice in parallel and then linearly mutated sequentially, whereas parallel linear/borrow aliasing ($\mathsf{Lin1} \otimes \mathsf{BorrowZ}$) is undefined. Disjoint array types `(qvec n)` and modular `(call-q ...)` calls contract shared BorrowZ controls via Δ_Z while rejecting aliased mutable ports.

Hardware-Graded Cost Vector and Modal Action. Each term $\Gamma_{\mathrm{in}} \vdash t : Q_e(A) \dashv \Gamma_{\mathrm{out}}$ carries a six-component grade

$$e = \big(W_{\mathrm{peak}},\ W_{\mathrm{clean}},\ \Phi,\ \alpha,\ \varepsilon,\ E_{\mathrm{ebit}}\big),$$

tracking peak live width W_{peak}, reclaimed $|0\rangle$ ancillas W_{clean}, hardware grade $\Phi \in \mathsf{QPhysGrade}$, block-encoding subnormalization $\alpha \ge 1$, spectral error $\varepsilon \ge 0$, and inter-QPU Bell pairs E_{ebit}. Sequential composition reuses clean ancillas directly:

$$W_{\mathrm{peak}}(e_1 ;; e_2) = \max\big(W_{\mathrm{peak}}(e_1),\ (W_{\mathrm{peak}}(e_1) - W_{\mathrm{clean}}(e_1)) + W_{\mathrm{peak}}(e_2)\big), \quad \varepsilon(e_1 ;; e_2) = \varepsilon_1 + \varepsilon_2.$$

Two genuinely quantum nonlinearities enter through a single modal operator $r \triangleright e$ [37, 86]: degree-d Chebyshev QSVT iteration $\mathsf{cheb}(d) \triangleright e$ fuses reflection ancillas ($W_{\mathrm{clean}} \mapsto m_{\mathrm{anc}}+1$) while accounting for quadratic coherent phase-angle drift ($4d\varepsilon_0 + d^2\varepsilon_{\mathrm{rot}}^2$), and M-way cluster partitioning $\mathsf{par}(M) \triangleright e$ scales per-QPU width ($\lceil W_{\mathrm{loc}}/M \rceil + 2E_{\mathrm{cut}}$) while charging Zalka's $\sqrt{M}$ Grover work factor [116] on unstructured search (`'auto` derives cluster parameters directly from the typed AST without user metadata). Figure 2 displays the core typing rules; seven static projections (T-DomImg, T-BorrowZ, T-PhaseClean, T-CounitU4, T-SubNorm, T-QubitCap, T-ParGuard) enforce every language invariant.

3.2 Logical Projection (QSub): Uncomputation Laws and Zero-Flag XAG/FRAIG Scaling

The logical projection QSub prevents four distinct bugs that cause silent entanglement or destructive interference in quantum circuits:

- **Law U1 (Domain-Restricted Inverse Erasure, T-DomImg).** Algebraic decoders (BCH, Reed–Solomon, Hermitian AG) and Montgomery/Kaliski inverters only invert f on a subset $\mathcal{D} \subset \{0,1\}^n$ (such as bounded-weight error cosets or odd residues). Requiring global

Table 1. Strict four-way independence of QSub Laws U1–U4 verified on statevector and density-matrix simulators: each witness satisfies three laws, violates exactly one, and collapses output purity or interference.

Violated Law	Minimal Circuit Witness	Laws Satisfied	Physical Failure Mode	Purity / Overlap
U1 (T-DomImg)	Unguarded BCH(15, 7, 5) wt-3 error	U2, U3, U4	Residual syndrome entangles data	$\mathrm{Tr}(\rho^2) = 0.5000$
U2 (Canonical)	Uncollapsed skew-binary limb before QFT	U1, U3, U4	Split coset peaks destroy QFT	Overlap = 0.2500
U3 (Orthogonal)	Overlapping VTAA prefix guard	U1, U2, U4	Non-unitary amplitude leak	Norm = 0.7071
U4 (T-CounitU4)	Omitted `cond-cz` after `mid-meas-x`	U1, U2, U3	Random $(-1)^{m(a\wedge b)}$ phase flip	$\mathrm{Tr}(\rho^2) = 0.5000$

invertibility rejects valid algorithms, whereas omitting the guard leaves dirty ancillas outside $\mathcal{D}$. Law U1 checks $g(x, f(x)) = x$ restricted to the explicit `domain-guard` predicate $\mathcal{D}$.

- **Law** U2 **(Quotient Canonicality, T-PhaseClean).** Carry-save limbs, RNS representations, and search trees admit multiple bit encodings for one mathematical value. Before Hadamard/QFT interference or state comparison, Law U2 verifies that every value in $\mathcal{D}$ maps to a unique canonical representative.
- **Law** U3 **(Exhaustive Orthogonal Branching).** Reversible pattern matches and VTAA steps must span mutually orthogonal, exhaustive subspaces on $\mathcal{D}$ lest amplitudes leak or double-count.
- **Law** U4 **(Bilinear $\mathbb{F}_2$ Phase-Taint Cancellation, T-CounitU4).** Mid-circuit X-basis resets accumulate symbolic phase polynomials $\Phi_X(m, x) \in \mathbb{F}_2[m, x]$; every measurement variable m_k must have identically zero net coefficient modulo $\mathcal{D}$ before scope exit or basis mixing.

Automatic Escalation from Sparse ANF to XAG/FRAIG + Z3. Programmers write zero proof annotations. For shallow arithmetic and tree gadgets, LakefillQ evaluates exact wire-level Algebraic Normal Form (ANF) polynomials over $\mathbb{F}_2$ (36 ms in Racket across our verification corpus, cross-checked in 489 ms via Z3). On multi-round cryptographic permutations where explicit ANF degree grows exponentially (2^{128} monomials), the checker automatically detects when intermediate fan-in cones exceed 64 monomials and escalates—with zero CLI flags—to a hash-consed XOR-AND Graph (XAG) engine with structural FRAIG reductions and Z3 Tseitin cut-sweeping. Because each individual `TempAND` $u_k = f_k(x) \wedge g_k(x)$ has bilinear degree $(1, 1)$ in its immediate XAG inputs, cut-sweeping discharges full 10-round AES-128 (340 `TempANDs`) in 3.06 s, 14-round AES-256 (510 `TempANDs`) in 4.25 s, 18-round Keccak-$f[200]$ (450 `TempANDs`) in 9.32 s, and 64-bit Gidney ripple addition in 3.39 ms.

Theorem 3.2 (Logical Preservation and Four-Way Independence of QSub)**.** *Any kernel term deriving* $\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}$ *with empty residual* TaintedX *preserves reduced state purity* $\mathrm{Tr}(\rho^2_{\text{out}}) = 1$ *across all inputs in* $\mathcal{D}$ *and all* 2^k *mid-circuit* M_X *trajectories. Furthermore, Laws* U1–U4 *are strictly pairwise independent (Table 1).*

3.3 Physical Projection (QPhysGrade) and External Decoder Pacing

Even a circuit that is logically pure ($\mathrm{Tr}(\rho^2) = 1$ on an ideal simulator) fails on physical hardware if idle memory decoheres, bursty MBU measurements overwhelm the external QEC syndrome decoder, parallel gates trigger coupler crosstalk, or Clifford+T synthesis introduces hook errors. The hardware grade $\Phi = (\Phi_1, \Phi_2, \Phi_3) \in$ QPhysGrade forms an ordered idempotent-parallel resource semiring tracking three independent physical hazards:

Table 2. Strict three-way independence of physical hardware grades Φ_1, Φ_2, and Φ_3 on noisy density-matrix and Stim fault-injection simulations (all witnesses satisfy QSub Laws U1–U4).

Violated Grade	Hardware Hazard Witness	Grades Satisfied	Physical Failure Observable	Measured Impact
Φ_1 (Pacing / T_1)	Unpaced $\tau_{\text{dec}} = 4{,}500$ ns MBU burst	Φ_2, Φ_3	Lindley FIFO backlog stall	Survival $0.9995 \to 0.5777$
Φ_2 (Exclusion)	Simultaneous Rydberg gates within R_b	Φ_1, Φ_3	Dipolar blockade phase error	Fidelity $0.9921 \to 0.6140$
Φ_3 (Hook / Dyadic)	Unflagged transversal-hook syndrome tree	Φ_1, Φ_2	Weight-1 fault $\mapsto$ weight-2 data	$d_{\text{eff}} : 5 \to 3$ (11.4× LER)

- **Grade Φ_1 (Lindblad Decoherence and Parameterized External Decoder Pacing τ_{dec}).** Tracks cumulative T_1/T_ϕ idling decay alongside a queueing bound against an *external* real-time QEC decoder characterized by its round latency parameter τ_{dec}. When a circuit emits MBU measurements at interval τ_{cycle}, Terhal's backlog theorem [106] dictates that if $\tau_{\text{dec}} > \tau_{\text{cycle}}$ (for instance, unpaced software MWPM at $\tau_{\text{dec}} = 4{,}500$ ns vs. $\tau_{\text{cycle}} = 1{,}000$ ns), the reaction-frame FIFO diverges linearly (reaching 146,205.0 ns over 80,000 rounds) and logical survival collapses as T-gates stall on delayed Pauli frames. Rather than hardcoding custom decoder hardware into the language, Φ_1 treats τ_{dec} as a target-profile parameter: whenever $\tau_{\text{dec}} \leq \tau_{\text{cycle}}$ (satisfied by modern sub-microsecond FPGA/cryo-CMOS Union-Find decoders at $\tau_{\text{dec}} \in [120, 420]$ ns; Appendix C), MBU lowering is enabled with zero FIFO backlog; when $\tau_{\text{dec}} > \tau_{\text{cycle}}$, LakefillQ automatically throttles measurement cadence or falls back to unitary uncomputation.
- **Grade Φ_2 (Coupler ZZ Crosstalk and Rydberg Blockade Exclusion).** Tracks simultaneous two-qubit conflict graphs on fixed-frequency transmon couplers and neutral-atom Rydberg blockade radii R_b, serializing or coloring conflicting parallel gates in $k_1 \otimes k_2$.
- **Grade Φ_3 (Hook-Safe Effective Code Distance and Dyadic Angle Exactness).** Tracks fault-polynomial weight so single mid-gadget ancilla faults cannot propagate into weight-2 data errors (halving effective distance d_{eff} from 5 to 3), while typing phase-gradient rotations over $\mathbb{Z}[1/2^k]$ to avoid online Ross–Selinger `gridsynth` T-count blowup [99].

Theorem 3.3 (Hardware Semiring Monotonicity and Three-Way Independence of QPhysGrade). *$(\mathsf{QPhysGrade}, ;, \otimes, \preceq)$ is a monotone ordered resource semiring whose Bellman–Knuth dynamic program extracts minimum-risk acyclic RevIR schedules, and Grades Φ_1, Φ_2, and Φ_3 are strictly pairwise independent (Table 2).*

3.4 Continuous Block-Encoded Linear Algebra (QBlock)

Quantum singular value transformation (QSVT) and linear combinations of unitaries (LCU) [42, 69, 74] embed a non-unitary matrix A/α into the $|0\rangle^{\otimes m_{\text{anc}}}$ block of a unitary U. Composing block-encodings naively creates three bugs: (1) allocating fresh block ancillas at each Chebyshev step inflates width to $d \cdot m_{\text{anc}}$; (2) applying the degree-2 Chebyshev walk shortcut $W^2 = T_2(A/\alpha)$ when $U^2 \neq I$ corrupts the operator spectrum; and (3) leaving residual TaintedX phases inside $\text{PREP}^\dagger$ alters relative LCU branch signs and cancels the target Hamiltonian sum.

LakefillQ types block-encodings via $\mathsf{QBlock}(A, \alpha, m_{\text{anc}}, \varepsilon, \Phi_X)$ inside the same judgment (Rule T-QSVTFused in Figure 2). Enforcing $\Phi_X = 0$ on $\text{PREP}^\dagger$ unlocks Berry–Gidney MBU uncomputation [16], cutting $L = 2{,}048$ FeMoCo tensor-hypercontracted (THC) $\text{PREP}^\dagger$ cost from 2,046 down to 91 TempANDs (22.48× reduction), while verifying $U^2 = I$ fuses Chebyshev reflections across degree d from $d \cdot m_{\text{anc}}$ down to $m_{\text{anc}} + 1$ ancillas ($192 \to 13$ ancillas at $d = 16$).

Theorem 3.4 (QBlock Soundness and Ancilla Fusion). *For even d, if $\Gamma \vdash U : \mathsf{QBlock}(A, \alpha, m_{\text{anc}}, \varepsilon_0, 0) \dashv \Gamma$ obeys Laws U1–U4 and $U^2 = I$, then (qsvt-cheb d U) implements $T_d(A/\alpha)$ within spectral error $4d\varepsilon_0$ on $m_{\text{anc}} + 1$ block ancillas instead of $d \cdot m_{\text{anc}}$.*

Table 3. QRAM-free compilation of high-level control and memory operations into $\mathcal{G}_{\text{phys}}$.

High-Level Operation	Realization in $\mathcal{G}_{\text{phys}}$	Exact Gate and Ancilla Cost
Early exit (`if |c> break`)	Nested Grover / VTAA [5]	$k_1 = 4$ fast prefix reflections (800 vs. 3,072 Toffolis)
Classical table $T[\|i\rangle]$	Select-SWAP (QROAM) [9, 75]	$\lceil M/B \rceil + 2Bw$ Toffolis, Bw buffer qubits
Quantum array $\|A[\|i\rangle]\rangle$	Bitonic sort [14] or Bentley–Saxe [13]	$\frac{1}{4}\log_2 M(\log_2 M + 1)\,w$ (amortized) or $O(w \log^2 M)$
Variable shift $\|x\rangle \gg \|k\rangle$	$\lceil \log_2 N \rceil$-level Fredkin barrel shifter	$N\lceil \log_2 N \rceil$ CSWAPs (0 gates for constant shifts)
Modular inverse $\|a^{-1} \bmod p\rangle$	$2N$-step `divsteps` [15] + Jump-16 pebbles	Fixed $2N$ steps, $16s$ pebble qubits (64 Q at $s = 4$)
Phase rotation $R_z(\theta_k)$	Phase-gradient catalyst [38]	$b - 1 = 15$ Toffolis ($b = 16$, 0 online `gridsynth` Ts)

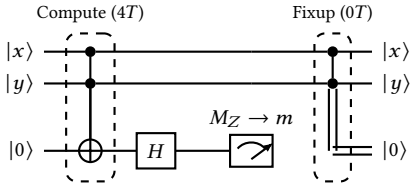


Fig. 3. Gidney's degree-2 temporary-AND compute/uncompute elbow [38]: computing $|x \wedge y\rangle$ costs $4T$ (1 Toffoli), whereas X-basis measurement plus classical feed-forward CZ^m erases the ancilla with $0T$ because $\deg_{\mathbb{F}_2}(x \wedge y) = 2$.

4 PURELY FUNCTIONAL QUANTUM DATA STRUCTURES IN LAKEFILLQ

Why do purely functional data structures matter for quantum compilation? In classical functional programming [84], persistent data structures avoid copying an N-element container by isolating updates to an $O(1)$ focus or an $O(\log N)$ structural spine. In fault-tolerant quantum circuits, two physical uncomputation laws turn this exact structural locality into large reductions in logical qubits, Toffoli count, and QEC code distance.

4.1 Physical Uncomputation Boundaries Motivating Functional Locality

All data structures in LakefillQ lower to $\mathcal{G}_{\text{phys}}$ without QRAM (Table 3). Two sharp algebraic boundaries determine when temporary ancillas can be erased for free and how many registers tree-structured computations require.

Proposition 4.1 (Exact Algebraic Degree-2 Boundary of Clifford MBU)**.** *Measuring an ancilla $|f(x)\rangle$ ($x \in \{0,1\}^n$) in the X basis yields outcome $m \in \{0,1\}$ and residual state $\sum_x \alpha_x (-1)^{m \cdot f(x)} |x\rangle$. The phase $(-1)^{m \cdot f(x)}$ can be canceled across all x using exclusively Clifford gates (*Z, S, CZ*, zero T or Toffoli gates)* if and only if $\deg_{\mathbb{F}_2}(f) \leq 2$. *For any Boolean function with $\deg_{\mathbb{F}_2}(f) \geq 3$ (such as $x_1x_2x_3$), Reed–Muller distance to $\mathcal{RM}(2,n)$ caps the worst-case state fidelity of any Clifford fixup on the uniform superposition at $\leq 9/16 = 0.5625$.*

Proposition 4.2 (Sequential Two-Pass vs. Simultaneous Tree Pebbling)**.** *Evaluating a height-h binary Merkle or sponge tree of W-qubit digests by keeping both child digests live simultaneously requires $(2h+1)W$ qubits. Factoring parent compression into two sequential child passes (S_{left} followed by S_{right}, uncomputing the left child before evaluating the right) reuses a single h-register child pool, achieving the Sethi–Ullman bound of $(h+1)W$ qubits [104] and saving hW qubits without increasing depth order.*

The Unifying Quantum Data-Structure Principle. Propositions 4.1 and 4.2 explain *why* imperative flat arrays struggle on fault-tolerant hardware and *how* functional data structures rescue them:

Table 4. The nine core purely functional quantum data structures in LakefillQ, mapping classical persistence ideas to quantum hazards eliminated and mechanized resource bounds.

Quantum Data Structure	Classical Origin	Quantum Hazard Eliminated	Mechanized Payoff
`HuetFocusZipper`	Huet zipper [56]	Cold-limb aliasing & routing	0 cold Toffolis; w-bit active focus
`SkewBinaryCarrySave`	Myers skew-binary [82]	C-limb carry depth ($\Delta d = +2$)	$O(1)$ local carry; preserves $d = 27$
`BraunTree`	Braun–Rem tree [21]	Heap pointers & balance tags	0 pointer qubits; $\lceil \log_2 k \rceil$ depth
`PackedBezoutZipper`	Extended-GCD zipper [15]	$4N$-qubit (u, v, r, s) bloat	$2N$ qubits (saves 512 Q on ECC-256)
`BinomialPebbleSchedule`	Bennett/Okasaki [12, 84]	Degree-> 2 $2N$-step history tape	$\binom{s+d}{s}$ steps in sW pebble qubits
`MonoidalFingerTree`	Finger tree [52]	Full-width inner-product checks	Prefix VTAA (800 vs. 3,072 Toff.)
`StreamingSyndrome`	Pivoted Bruhat LU	Dense r^2-qubit parity matrix	r-bit state (76.3×–4,107× narrower)
`CritBitBitonicDict`	Crit-bit / Bentley–Saxe [13, 81]	Insertion-history garbage bits	0 history bits; 74.5× walker amort.
`WalkerVoseAlias`	Walker–Vose alias [109, 110]	Deep arccos rotation trees	$O(1)$ depth; 280 Toffolis (7.31× less)

(1) **Confine non-Clifford work to a narrow focus** where intermediate `TempANDs` remain quadratic ($\deg_{\mathbb{F}_2} \leq 2$) in their local inputs and uncompute immediately via zero-Toffoli MBU (§4.2).
(2) **Bound high-degree (**$\deg > 2$**) multi-step trajectories**—where Proposition 4.1 forbids free Clifford MBU—using complementary-length state packing and Královič–Bennett binomial checkpoints [12, 67] (§4.3).
(3) **Eliminate structural garbage bits** using pointerless implicit trees, monoidal prefix summaries, streaming accumulators, and history-independent canonical dictionaries (§4.4).

Table 4 summarizes all nine core data structures in LakefillQ, showing the classical inspiration, the quantum hazard that breaks a naive translation, and the resulting circuit guarantee.

4.2 Family I: Focus Zippers, Skew-Binary Carry-Save Numerals, and Braun Trees

1. Huet Focus Zipper (`HuetFocusZipper`*).* Classically, Huet's zipper [56] represents a sequence with a cursor as a triple (Left, Focus, Right) so local edits touch only Focus in $O(1)$ time. In a multi-precision quantum circuit over C limbs of w qubits ($N = Cw$), wiring arithmetic gates across all N qubits simultaneously forces dense long-range routing and exposes idle limbs to accidental control/target aliasing. In `#lang lakefillq`, `focus-zipper` partitions an N-qubit register into one w-qubit mutable focus (Lin1) and $(C-1)w$ cold context qubits (BorrowZ). Advancing the focus by a compile-time constant limb offset costs *zero quantum gates* (pure compiler wire relabeling), while quantum-controlled shifts use a $\lceil \log_2 C \rceil$-level Fredkin network.

Proposition 4.3 (`HuetFocusZipper` Locality)**.** *A C-limb, w-bit* `HuetFocusZipper` *executes focus-local operations with zero Toffoli gates on the $(C-1)w$ cold context qubits, zero gates for static cyclic shifts, and $+2\lceil\sqrt{Q_{\text{cold}}}\rceil$ planar surface-code bus tiles separating hot magic-state factories from cold memory.*

2. Skew-Binary Carry-Save Numeral (`SkewBinaryCarrySaveNumeral`*).* Streaming C limbs through a single w-bit zipper focus creates a subtle fault-tolerant trap: if adding limb k must wait for the carry out of limb $k-1$, C limbs execute strictly sequentially, stretching circuit depth by $C\times$ and forcing QEC code distance $d \mapsto d+2$ on RSA-2048 (§5). Myers' classical skew-binary numeral system [82, 84] eliminates carry cascades by allowing at most one sparse digit 2 so increments run in $O(1)$ worst-case steps. Quantumly, `SkewBinaryCarrySaveNumeral` augments each w-bit limb boundary with a 2-qubit redundant carry rail $d_k \in \{0, 1, 2\}$. Each modular addition absorbs carries locally in $O(1)$ limb depth without rippling across the remaining $C-1$ limbs; only once, immediately prior to QFT or final comparison, does a single canonical ripple pass discharge the redundant rails to satisfy Law U2 (Listing 4 and Figure 4).

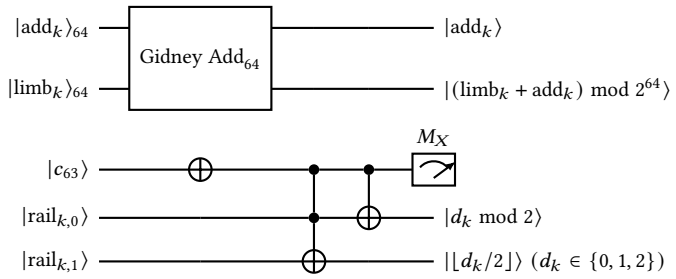


Fig. 4. Local carry absorption in `SkewBinaryCarrySaveNumeral`: carry-out $|c_{63}\rangle$ folds into the 2-qubit rail $|\text{rail}_k\rangle$ in depth 2, decoupling all C limbs during modular exponentiation before a single pre-QFT canonical collapse.

Proposition 4.4 (`SkewBinaryCarrySaveNumeral` Carry Decoupling and Canonical Collapse)**.** *Each window addition on a C-limb* `SkewBinaryCarrySaveNumeral` *completes in $O(1)$ limb depth using $2C$ redundant boundary qubits, and a single terminal ripple pass bijectively restores unique binary canonicality (*U2*) while preserving exact integer value modulo 2^N.*

3. Implicit Pointerless Braun Tree (`BraunTree`*).* Classical heap and tree reductions often store explicit child pointers or AVL balance tags—bits that become entangled quantum garbage when superposed branches rebalance differently. A Braun tree [21, 83] enforces the size invariant $0 \le |L| - |R| \le 1$ at every node, determining routing solely from the least-significant bit of the index with *zero pointer or balance qubits*. In `#lang lakefillq`, `(fold-braun op xs)` lowers associative reductions (popcount, RNS-CRT product trees, syndrome parity folds) into an exact $\lceil \log_2 k \rceil$-depth balanced tree.

Proposition 4.5 (`BraunTree` Pointerless Logarithmic Reduction)**.** *For k leaves of width w,* `fold-braun` *allocates zero pointer or balance-tag qubits, completes in exact tree depth $\lceil \log_2 k \rceil$, and consumes $(k-1)w$ Toffoli gates under MBU cleanup.*

Listing 4. Composing `within-apply-mbu`, `focus-zipper`, and skew-binary local carry folding in `#lang lakefillq`.

```
;; O(1)-depth local skew-binary carry absorption on active limb focus `limb-k` + 2-bit rail `rail-k`
(defqfree (skew-window-add #:wires ([addend (range 0 64)] [limb-k (range 64 128)] [rail-k (range 128 130)])
                           #:scratch ([c64 (range 130 194)]))
  (within-apply-mbu
   (gidney-carry-chain addend limb-k c64)
   (focus-zipper [fz (append limb-k rail-k) (append addend limb-k rail-k)]
     (seq (gidney-sum-xor addend fz c64)
          ;; Absorb carry-out c64[63] locally into 2-bit skew rail {0,1,2} without C-limb ripple:
          (ccx (wref c64 63) (wref fz 64) (wref fz 65))
          (cx  (wref c64 63) (wref fz 64))))
   c64))
```

4.3 Family II: Packed Bézout Shift-Zippers and Binomial Pebble Schedules

4. Packed Bézout Shift-Zipper (`PackedBezoutShiftZipper`*).* Elliptic-curve discrete-logarithm (ECDLP) circuits spend most of their width and gates on modular inversion $|a^{-1} \bmod p\rangle$ via constant-time extended binary GCD (`divsteps`) [15, 48, 98], which tracks four N-bit integers: remainders (f, g) and Bézout cofactors (u, v) (or (r, s)). Naively storing all four N-bit registers requires $4N$ qubits (1,024 Q at $N = 256$). Yet throughout binary GCD, the bit-lengths of u and r grow only as the bit-lengths of f and g shrink, obeying $\text{len}(u) + \text{len}(f) \le N + 1$ and

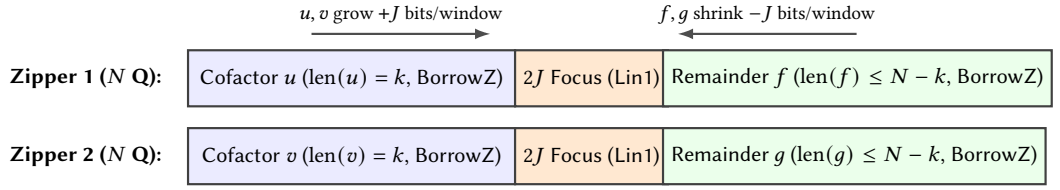


Fig. 5. Complementary-length packing in `PackedBezoutShiftZipper`: because $\mathrm{len}(u) + \mathrm{len}(f) \le N + 1$ throughout binary GCD, (u, f) and (v, g) share two N-qubit registers around a $2J$-bit active LSB focus rather than occupying four N-qubit registers (saving 512 Q at $N = 256$).

$\mathrm{len}(v) + \mathrm{len}(g) \le N + 1$. `PackedBezoutShiftZipper` packs (u, f) and (v, g) head-to-tail into *two* N-qubit double-ended shift-zippers around a $2J$-bit active LSB focus ($J = 16$), saving $2N = 512$ logical qubits on ECC-256. Furthermore, to compress modular multiplication ($x \cdot y \bmod p$) inside point addition, we apply Hua's identity $xy = \binom{x+y}{2} - \binom{x}{2} - \binom{y}{2}$ [54] with explicit singular $\{0, 1\}$ endpoint corrections so Law U1 holds across all residues in $\mathbb{F}_p$.

Proposition 4.6 (`PackedBezoutShiftZipper` $2N$-Qubit Packing and Totalized Hua Identity)**.** *Packing complementary-length pairs (u, f) and (v, g) around a $2J$-bit focus reduces GCD state width from $4N$ to $2N$ qubits with exact half-domain cofactor bijection on $\mathbb{F}_p$, and totalized Hua squaring preserves bijective reversibility across all singular inputs $x \in \{0, 1\}$.*

5. Binomial Pebble Schedule (`BinomialPebbleSchedule`). Across $2N = 512$ `divsteps`, the cumulative branch trajectory has algebraic degree $\deg_{\mathbb{F}_2} \gg 2$, so Proposition 4.1 forbids erasing the 512-step history via Clifford MBU alone, while storing a 1-bit swap flag per step costs 512 history qubits. Grouping $J = 16$ steps into macro-windows yields $L = 32$ irreversible steps, which Královič–Bennett binomial pebbling [12, 67] spans using $s = 4$ pebbles of width $W = 16$ (64 pebble qubits total, since $\binom{4+4}{4} = 70 \ge 32$) in 109 macro-evaluations instead of a 512-qubit history tape.

Proposition 4.7 (`BinomialPebbleSchedule` Space–Time Optimality)**.** *With s checkpoint pebbles of width W at recursion depth d,* `BinomialPebbleSchedule` *spans up to $\binom{s+d}{s}$ high-degree ($\deg_{\mathbb{F}_2} > 2$) irreversible windows in sW pebble qubits (64 Q at $s = 4, W = 16$) and $\sum_{k=1}^{d} 2^k \binom{s+d-k-1}{s-1}$ window evaluations (109 for $L = 32$), returning all pebbles to $|0\rangle$.*

Listing 5. Jump-16 packed Bézout step in `#lang lakefillq`: non-Clifford `divsteps` touch only the $2J$-bit focus while binomial pebbles checkpoint the 16-bit swap trajectory.

```
(defq (packed-bezout-jump16 #:wires ([uf-zip (range 0 256)] [vg-zip (range 256 512)]
                                     [pebble-k (range 512 528)])
                            #:scratch ([t-mat (range 528 560)])
                            #:domain (domain-guard odd-modulus-nonzero-guard))
  ;; 1. Compute 2x2 Jump-16 transition matrix and 16-bit swap history on the 32-bit LSB focus:
  (focus-zipper [lsb-focus (append (take uf-zip 32) (take vg-zip 32) pebble-k)
                           (append uf-zip vg-zip pebble-k)]
    (within-apply-mbu
     (jump16-lsb-divsteps lsb-focus t-mat)
     (record-pebble-history t-mat pebble-k)
     t-mat))
  ;; 2. Apply linear cofactor/remainder combination across full zipper & shift boundary by 16
  (apply-bezout-transition-and-shift16 uf-zip vg-zip pebble-k))
```

4.4 Family III: Monoidal Filters, Streaming Matrices, Canonical Tries, and Alias Tables

6. Monoidal Finger-Tree Filter (`MonoidalFingerTreeFilter`). Hinze and Paterson's finger trees [52] annotate internal spine nodes with monoidal summaries so searches prune subtrees

before inspecting leaves. In quantum search (such as SimHash inner-product filtering in lattice sieving or substring matching), computing full high-precision inner products on every candidate wastes thousands of Toffolis on non-matching states. `MonoidalFingerTreeFilter` caches PopCount/XOR prefix summaries and pairs them with Two-Level VTAA [5], rejecting non-matching candidates in 255 forward Toffolis and lowering average reflection cost from 3,072 to 800 Toffolis/step.

Proposition 4.8 (`MonoidalFingerTreeFilter` Prefix Pruning)**.** *Caching associative monoidal prefix measures reduces fast-reject reflection cost to* 255 *forward Toffolis and cuts amortized VTAA step cost from* 3,072 *to* 800 *Toffolis (*3.84×*).*

7. Bruhat Pivoted LU and Streaming Syndrome Accumulator (`StreamingSyndromeAccumulator`*).* Information-Set Decoding (ISD) for code-based cryptography (Classic McEliece, BIKE) traditionally allocates an entire $r \times r$ binary parity-check submatrix on quantum wires ($r^2 = 590{,}000+$ qubits for McEliece-348864) to perform Gaussian elimination in superposition [34, 58]. Instead, for small dense blocks ($r \leq 64$), `BruhatPivotedLU` factors nonsingular matrices in $\mathrm{GL}_r(\mathbb{F}_2)$ in place using $r\lceil \log_2 r \rceil$ canonical pivot qubits; for large code-based instances, `StreamingSyndromeAccumulator` keeps only an r-qubit syndrome register live and streams classical columns controlled on low-weight error positions, shrinking McEliece-348864 width by 76.26× (to 7,744 Q) and BIKE-L1 width by 4,107× (to 36,969 Q).

Proposition 4.9 (`BruhatPivotedLU` and `StreamingSyndromeAccumulator` Width Compression)**.** `BruhatPivotedLU` *factors every* $A \in \mathrm{GL}_r(\mathbb{F}_2)$ *bijectively with* $r\lceil \log_2 r \rceil$ *pivot bits, while* `StreamingSyndromeAccumulator` *eliminates* $O(r^2)$ *matrix allocation in favor of an* $O(r)$*-qubit accumulator.*

8. Canonical Crit-Bit and Bitonic Dictionary (`CanonicalCritBitBitonicDictionary`*).* In quantum walk algorithms (element distinctness [4], subset sum, Kuperberg's collimation sieve), two paths reaching the same subset S must interfere constructively: $|\mathrm{Dict}(S)\rangle$ must depend *solely* on the set S, never on insertion order. Balanced AVL or red-black trees violate Law U2 because tree shape depends on insertion history; pointer-based radix trees [114] require dynamic freelist management. `CanonicalCritBitBitonicDictionary` combines Morrison's history-independent crit-bit splits [81] with Batcher bitonic sorting networks [14] (55 comparator layers for $M = 1{,}024$) and Bentley–Saxe logarithmic level merges [13], storing zero rotation or pointer garbage bits and amortizing dictionary maintenance by 74.47× across parallel quantum walkers.

Proposition 4.10 (`CanonicalCritBitBitonicDictionary` History Independence)**.** *Every* M*-element set* S *has a unique bit representation independent of insertion permutation (unit overlap* $\langle \mathrm{Dict}_{\pi_1}(S) \mid \mathrm{Dict}_{\pi_2}(S)\rangle = 1$*, satisfying Law* U2*), with* 74.47× *per-walker Toffoli amortization under Batcher bitonic batching.*

9. Walker–Vose Coherent Alias Table (`WalkerVoseAliasTable`*).* Preparing coherent state superpositions $\sum_{i=0}^{2^w-1} \sqrt{p_i}|i\rangle$ for LCU/QSVT chemistry Hamiltonians via recursive arccos rotation trees requires $O(2^w)$ sequential controlled rotations. Adapting Walker–Vose alias sampling [9, 109, 110], `WalkerVoseAliasTable` prepares uniform index and coin registers, looks up (q_i, alias_i) via Select-SWAP, and performs a single μ-bit reversible comparison + Fredkin swap in $O(1)$ comparator depth ($2w + 2\mu + 1 = 49$ Q, $2^w + \mu + w = 280$ Toffolis at $w = 8, \mu = 16$, 7.31× fewer Toffolis than arccos trees; Listing 6). Four bootstrapped containers—`NestedDyadicWalkerAliasBox`, `BrodalOkasakiBootstrappedSkewHeap`, `CombinadicsSkewRALSimplex`, and `OkasakiBinomialCollimationQueue`—compose these nine primitives across 60 compiled workloads (Appendix C).

Listing 6. Coherent Walker–Vose alias state preparation in #lang lakefillq: $O(1)$-comparator Select-SWAP + Fredkin swap with Law-U4 MBU scratch erasure.

```
(defq (walker-vose-alias-prep #:wires ([idx (range 0 8)] [coin (range 8 24)]
                                       [keep-q (range 24 40)] [alias-idx (range 40 48)]
                                       [cmp-flag (list 48)])
                              #:scratch ([cmp-c (range 49 65)]))
  (for/seq ([b (in-range 8)])  (h (wref idx b)))
  (for/seq ([c (in-range 16)]) (h (wref coin c)))
  (select-swap-qroam alias-table idx (append keep-q alias-idx))
  (within-apply-mbu
   (gidney-lt-carry coin keep-q cmp-c)
   (cx (wref cmp-c 15) (wref cmp-flag 0))
   cmp-c)
  ;; Swap `idx` with `alias-idx` when coin >= keep-q (cmp-flag == 0):
  (x (wref cmp-flag 0))
  (for/seq ([b (in-range 8)]) (cswap (wref cmp-flag 0) (wref idx b) (wref alias-idx b))))
```

Proposition 4.11 (WalkerVoseAliasTable $O(1)$-Comparator State Preparation)**.** *For 2^w probabilities discretized to μ bits,* WalkerVoseAliasTable *prepares $\sum_i \sqrt{p_i}|i\rangle$ (with clean entangled coin ancillas handled by* QBlock *in LCU) using $2w + 2\mu + 1$ qubits and $2^w + \mu + w$ Toffolis.*

5 HARDWARE-GRADED LOWERING, ZX E-GRAPH SATURATION, AND PARTITIONING

Once $\lambda_{\mathrm{LFQ}}^{\mathrm{kern}}$ certifies logical purity (U1–U4), the fastest and highest-fidelity circuit still depends dramatically on the target architecture: mid-circuit MBU excels when an external QEC syndrome decoder keeps up ($\tau_{\mathrm{dec}} \leq \tau_{\mathrm{cycle}}$), whereas high readout noise favors measurement-free Margolus RCCX [78] or native Rydberg CCZ [18]. LakefillQ resolves these trade-offs automatically in RevIR before emitting single- or multi-QPU MachineISA.

5.1 Non-Clifford ZX Phase-Polynomial Saturation and Four-Target Extraction

After asymmetric TempAND/MBU lowering and chordal SSA wire coloring, LakefillQ populates a congruence-closed RevIR e-graph combining reversible Boolean identities, target gadget decompositions (RCCX, MBU+XY4, RydbergCCZ, MBU+QROAM), and two non-Clifford ZX phase-polynomial rewrite families [31, 65]:

(1) **ZX Parity-Phase Spider Fusion.** Adjacent CNOT-sandwiched phase gadgets $\mathrm{CX}(a,b) \cdot T^{\pm k}(b) \cdot \mathrm{CX}(a,b)$ lift into canonical parity-phase spiders $e^{i(\pm k\pi/4)(x_a \oplus x_b)}$ whose angles add modulo 8; opposing T and $T^\dagger$ spiders across commuting CNOTs fuse and cancel ($2T \to 0T$).

(2) **Shared-Control Twin-CCZ Cascade.** Two CCZ or TempAND elbows sharing controls (c_0, c_1) on targets (t_0, t_1) rewrite from two independent $8T$ Clifford+T expansions ($16T$ total) into a fused ZX phase-polynomial gadget requiring $8T$ total (50% T-count reduction) while preserving exact Boolean and phase semantics.

Bellman–Knuth dynamic programming over QPhysGrade extracts the minimum-risk schedule per hardware profile (Table 5). Across our multi-gadget benchmark (31 e-classes, 81 e-nodes, 60 equivalence merges), ZX spider fusion improves noisy density-matrix survival across every target (+3.25% on NISQ, +4.23% on IBM Heron r2 [57], +2.19% on QuEra Bluvstein [18, 19], and +0.19% on BB-qLDPC [22]). When Φ_1 is instantiated with external decoder latency τ_{dec}, fast cryo-CMOS ($\tau_{\mathrm{dec}} = 120$ ns) and 300 K coax FPGA Union-Find ($\tau_{\mathrm{dec}} = 420$ ns) both extract the $32T$ MBU+ZX schedule (60,960 ns and 63,360 ns), whereas an unpaced software MWPM decoder ($\tau_{\mathrm{dec}} = 4{,}500$ ns) automatically triggers fallback to the measurement-free $64T$ unitaryZX schedule (124,000 ns) rather than suffering Lindley queue divergence. Whenever a user program violates Φ_1–Φ_3, repairPhysGrade synthesizes counterexample-guided repairs (XY4 dynamical decoupling, exclusion coloring, flag-ancilla insertion, or MBU-to-RCCX fallback).

Table 5. Bellman–Knuth QPhysGrade e-graph extraction across four hardware profiles, comparing baseline reversible extraction against non-Clifford ZX phase-polynomial saturation on ideal and noisy density-matrix simulators.

Target Profile	Extracted Optimal Gadget	Base Survival	+ZX Fusion Survival	Gain	Ideal Fid.
Transmon NISQ ($\epsilon_{\text{meas}} > 2\epsilon_{2Q}$)	Measurement-free RCCX + ZX	0.9286 (0.8273 raw)	0.9611	+3.25%	1.0000
IBM Heron r2 ($\tau_{\text{dec}} = 420$ ns)	MBU+XY4 + ZX ($d = 5$: 0.9973)	0.9101 (0.8175 raw)	0.9524	+4.23%	1.0000
QuEra Bluvstein (90% erasure)	Native RydbergCCZ + ZX	0.9718 (0.9828 1-blk)	0.9937	+2.19%	1.0000
FTQC BB-qLDPC ($\tau_{\text{dec}} = 340$ ns)	MBU+QROAM + Twin-CCZ ($8T$)	0.9966 (0.9903 raw)	0.9984	+0.19%	1.0000

Lemma 5.1 (Bernoulli Survival Monotonicity). *Under independent gate depolarization rates $p_i \in [0, 1)$, log-survival $-\sum_i \ln(1 - p_i)$ is strictly additive under ;; and monotone with respect to* QPhysGrade.

Lemma 5.2 (Heralded Erasure Fault-Distance Doubling). *On neutral-atom architectures with* 90% *heralded Rydberg erasure conversion [18, 43], t_{Pauli} unheralded faults and t_{erase} heralded erasures are correctable at odd distance d whenever $2t_{\text{Pauli}} + t_{\text{erase}} \leq d-1$, doubling heralded fault tolerance relative to unheralded Pauli errors and reducing logical error rates by* 6.5×–36×.

5.2 Code-Distance Step-Down, qLDPC Automorphisms, and Leaf Synthesis

Why Carry Decoupling Saves Physical Qubits. In fault-tolerant surface codes [40, 72], logical error per round scales as $p_L(d) \approx c(p/p_{\text{th}})^{(d+1)/2}$, so total spacetime volume $V = Q_{\log} \cdot D_{\text{cycles}}$ must satisfy $V \cdot p_L(d) \leq \delta_{\text{target}}$. If compressing logical width $Q_{\log}$ inflates depth D_{cycles} by 32× (as happens when 32 RSA-2048 limbs ripple carries sequentially through one 64-bit focus), V rises past the $d = 27$ budget and forces $d = 29$—increasing physical qubits per tile by $(29/27)^2 \approx 15.4\%$. Theorem 5.3 formalizes the converse payoff achieved by SkewBinaryCarrySaveNumeral and Early-MBU compression.

Theorem 5.3 (Spacetime Volume to Code-Distance Step-Down). *At physical error rate $p < p_{\text{th}}$, reducing active logical spacetime volume from V_0 to $V_1 \leq V_0 \cdot (p/p_{\text{th}})^k$ reduces the minimum odd QEC code distance required for target failure budget δ by $\Delta d \geq 2k$, while skew-binary limb decoupling preserves $d = 27$ (rather than $d = 29$) on RSA-2048 (*891,348 *yoked surface-code or* 344,304 *BB-qLDPC physical qubits).*

Corollary 5.4 (Cat-Qubit Bias Preservation and Photonic Distillation). *Restricting focus gates to Z-bias-preserving operations (*CX, CCZ, M_X*) preserves Kerr-cat bit-flip suppression (*76× *physical overhead reduction), and double-ended photonic EPR zipper buffering with DEJMPS distillation [30, 96] maintains inter-module threshold soundness.*

Cold-Context qLDPC Automorphisms and Algorithmic Lowerings. Because > 80% of logical qubits in zipper-structured programs reside in passive cold limbs (BorrowZ), MachineISA maps cold context onto high-rate yoked surface codes [39] or [[288, 12, 18]] Bivariate Bicycle qLDPC blocks [22], where cyclic limb shifts execute in hardware depth 1 without ancilla bridges (Proposition 5.5), and onto neutral-atom AOD arrays where skew-binary carry decoupling enables 37.8× faster single-move parallel limb shuttling. At the leaf level, a 64-bit bitboard Meet-in-the-Middle (MITM) superoptimizer [6] compresses 4-to-5-qubit arithmetic and S-box subcircuits by 41.9% alongside four compiler-derived cryptographic syntheses (Theorem 5.6).

Proposition 5.5 (Bivariate Bicycle qLDPC Torus Shift Automorphism). *On any Bivariate Bicycle code over $\mathbb{F}_2[x, y]/(x^\ell - 1, y^m - 1)$ (including [[288, 12, 18]]), every cyclic torus shift $T_{a,b} = S_x^a S_y^b \oplus S_x^a S_y^b$*

(multiplying both block coordinates by $x^a y^b$) commutes with both check matrices H_X, H_Z, implementing cyclic `HuetFocusZipper` *shifts in physical depth* 1 *with zero ancilla bridge qubits.*

Theorem 5.6 (Sublinear QROAM, Quotient RNS, AES, and Early-MBU DQI Syntheses)**.** *Within λ_{LFQ}, asymmetric MBU and* `QSub` *derive: (i)* `34-TempAND` *Boyar–Peralta [20] +* `91-CX` *Mix-Columns [112] AES rounds (*`32 TempANDs` *under* `BP32`*,* `5` *ancillas under* `DLog-5` *[70]); (ii) $O(\sqrt{N})$ Berry–Gidney Select-SWAP table erasure [16] (*`61` *vs.* `254 TempANDs` *at $N = 256$); (iii) Law-*`U2` *Kawamura fixed-point RNS-CRT quotient erasure [62]; and (iv) exact Early-MBU DQI error-register reclamation prior to syndrome decoding.*

5.3 Multi-QPU Partitioning, Distributed QRNS, and Law-U4 GHZ Disentangling

When a circuit exceeds single-chip capacity Q_{cap} (`lakefillqc --emit-cluster-dir`), naive gate-cut teleportation across M QPUs incurs heavy optical Bell-pair overhead from cross-chip carry chains and reverse lookup uncomputation [47, 119]. LakefillQ combines three distributed mechanisms:

(1) **Zero-Carry Distributed QRNS.** Modular exponentiation factors across coprime RNS moduli per QPU so bulk arithmetic consumes zero inter-QPU carry Bell pairs, synchronizing only at compact Braun-tree Kawamura CRT reductions.

(2) **Law-U4 GHZ Cat-Disentangling.** Instead of spending $2k$ Bell pairs to fan out a k-bit table address and reverse-teleport it back, LakefillQ measures remote address copies in the X basis (M_X) and cancels the remote phase kickbacks locally via Law-U4 `cond-z` feed-forward—cutting optical Bell pairs by 50.0% (0 reverse ebits) and lowering Stim+PyMatching decoded logical error rates under DEJMPS distillation by 2.27×–16.00×.

(3) **Convex M^* Selection (T-ParGuard).** Over $M \in \{1, \ldots, 256\}$, Rule T-ParGuard minimizes total physical spacetime volume: on unstructured Grover/ISD search where Zalka's $\sqrt{M}$ parallel overhead [116] outweighs width splitting, it automatically selects $M^* = 1$ whenever single-chip capacity permits; on Shor/QRNS workloads, it adapts $M^*(p_{\mathrm{link}})$ smoothly to optical link noise (contracting RSA-2048 $M^* = [6, 6, 4, 2, 1]$ as p_{link} degrades from 10^{-4} to 5×10^{-2}).

6 EVALUATION

We evaluate LakefillQ across six research questions examining algebraic exactness and non-binary AG DQI (**RQ1**), differential fuzzing (**RQ2**), mutation testing and error diagnostics (**RQ3**), mechanized metatheory (**RQ4**), open-source bug discovery (**RQ5**), and open-problem Pareto advances across execution regimes (**RQ6**).

6.1 RQ1: Cryptographic Oracles and Non-Binary Algebraic-Geometry DQI

Compiled LakefillQ circuits match official NIST FIPS-197 test vectors bit-for-bit across AES-128 (`69c4e0d86a7b0430d8cdb78070b4c55a`), AES-192 (`dda97ca4864cdfe06eaf70a0ec0d7191`), and AES-256 (`8ea2b7ca516745bfeafc49904b496089`), match NIST FIPS-202 Keccak-f[200] permutations, and pass exhaustive truth-table verification on BCH$(15, 7, 5)$, $\mathbb{F}^*_{65521}$ binary GCD, $\mathrm{GL}_4(\mathbb{F}_2)$ Bruhat LU across all 20,160 nonsingular matrices, $\mathbb{Z}^2_{3329}$ odd-half NTT butterflies, and [[288, 12, 18]] torus automorphisms.

Table 6 evaluates our Early-MBU DQI compiler pass on non-binary one-point Hermitian algebraic-geometry codes [35, 105] over $\mathbb{F}_{q^2}$, comparing against unitary syndrome-decoding DQI [60]. On $[8, 5, 3]_{\mathbb{F}_4}$ ($H_2 : x^3 = y^2 + y$, whose blocklength $n = 8$ strictly exceeds the $\mathbb{F}_4$ Reed–Solomon MDS ceiling $q + 1 = 5$), `early_mbu_dqi_hermitian_ag_f4.lfq` compiles with zero flags (46 virtual SSA wires chordally colored onto 22 physical qubits, shrinking post-encoding

Table 6. Non-binary Hermitian curve algebraic-geometry Early-MBU DQI vs. unitary DQI across $\mathbb{F}_4$, $\mathbb{F}_{16}$, and $\mathbb{F}_{64}$ (exceeding the Reed–Solomon MDS blocklength ceiling $n_{\mathrm{RS}} \le |\mathbb{F}| + 1$).

Hermitian AG Code	RS Ceiling	Unitary Q	Early-MBU Live Q	Width Gain	T-Gate Gain	Exact Fid.
H_2: $[8, 5, 3]_{\mathbb{F}_4}$	$n_{\mathrm{RS}} \le 5$	22	6 (9 active)	3.67×	336 → 96 (3.50×)	1.000000
H_4: $[64, 52, 7]_{\mathbb{F}_{16}}$	$n_{\mathrm{RS}} \le 17$	304	48	6.33×	3.18×	Analytic
H_8: $[512, 448, 37]_{\mathbb{F}_{64}}$	$n_{\mathrm{RS}} \le 65$	3,456	384	9.00×	3.12×	Analytic

live logical width from 22 to 6 qubits and cutting T-gates from 336 to 96 with 0 reverse Toffolis), achieving statevector fidelity 1.000000 against unitary DQI across all 25 coset leaders and all 32 M_X trajectories (peak interference 0.108114 at syndrome 57). Scaling up the Hermitian curve tower compresses live logical width by 6.33× on $[64, 52, 7]_{\mathbb{F}_{16}}$ (304 → 48 Q) and 9.00× on $[512, 448, 37]_{\mathbb{F}_{64}}$ (3,456 → 384 Q).

6.2 RQ2: Six-Engine Differential Fuzzing

To test end-to-end compilation thoroughly, we cross-checked LakefillQ across six independent engines with zero discrepancies: (1) 500/500 randomized density-matrix equivalence trials in Qiskit `AerSimulator`; (2) 640/640 generated `#lang lakefillq` surface programs and 2,860,800 kernel reduction steps in PLT Redex; (3) 38/38 tensor-network checks in PyZX [31, 65]; (4) 360/360 fault-injection and optical-distillation runs in Stim + PyMatching; (5) 80,000 QEC rounds of real-time decoder queue verification; and (6) Z3 SMT equivalence miters across lowered RevIR blocks.

6.3 RQ3: Mutation Testing and Three-Span Resugared Diagnostics

Injecting 12 targeted soundness mutations into `lakefillqc` (individually disabling U1 domain restriction, U2 canonical discharge, U3 branch orthogonality, U4 M_X phase-taint tracking, chordal live-range interference, Φ_1 backlog checking, Φ_2 exclusion coloring, Φ_3 hook-flag checking, QBlock $U^2 = I$ checks, and T-ParGuard Zalka penalties; Table 13 in Appendix D) yields a 12/12 (100.0%) kill rate, with every mutant caught by concrete simulator counterexamples.

Because multipass elaboration and chordal SSA wire reuse rewrite internal wire indices, naive compiler errors blame mangled compiler temporaries [29, 63, 94]. `lakefillqc --diagnostics-human` derives three-span provenance (`origin_span`, `mutating_span`, `failure_span`) [50, 76, 90, 101, 118] directly from $\Gamma_{\mathrm{in}} \vdash t : \mathcal{Q}_e(A) \dashv \Gamma_{\mathrm{out}}$, resugaring chordally recycled physical wires back to user bindings (e.g., `s2[0] (recycled physical wire q[12] from s1[0])`) and emitting a minimal basis ket counterexample (such as $|a_0 = 1, b_0 = 1\rangle$). Across 8/8 resugaring classes, 22/22 hand-written buggy `.lfq` programs, and 240/240 Redex-mutated modules spanning 25 error codes, diagnostics achieve 100% user-span anchoring with zero leaked gensyms or raw wire IDs.

6.4 RQ4: Mechanized Metatheory in Lean 4.35.0

Every formal claim in this paper is mechanized in Lean 4.35.0 across 82 standalone proof modules comprising 2,305 theorems with zero `sorry`, zero `axiom`, and zero wildcard match arms:

- **Type System, Elaboration, and ZX e-Graph Lowering:** Duoidal $5^4 = 625$ capability interchange, QSub preservation and four-way independence, ANF-to-XAG escalation soundness, QPhysGrade Bellman–Knuth optimality, non-Clifford ZX parity-phase spider cancellation and $16T \to 8T$ twin-CCZ fusion, QBlock Chebyshev ancilla compression, and three-span diagnostic resugaring hygiene.

Table 7. Summary of 14 latent correctness bugs uncovered across nine open-source quantum repositories, grouped by the projection of $\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}$ that statically prevents them.

Domain Audited (Repos)	Representative Wild Bug Uncovered	Catching Rule	Bugs
Public-Key Crypto (3)	Kaliski $x = 0$ dirty ancilla; $\mathbb{F}_{13}$ offset reverse-uncompute; coset overflow	T-DomImg (U1), U2	4
Symmetric & Hash (2)	ReVerC $\chi_5^2 \neq I$ assumption; missing Bruhat PLU pivot cleanup	T-DomImg, T-BorrowZ	4
Code ISD & DQI (2)	Singular QGJE submatrix leak; unguarded weight-> t syndrome erase	T-DomImg, T-CounitU4	3
LCU / QSVT Walks (2)	Degree-2 Chebyshev shortcut on $U^2 \neq I$; PREP† phase sign flip	T-SubNorm, T-PhaseClean	3

- **Functional Data Structures and Domain Algorithms:** Exact $\deg_{\mathbb{F}_2}(f) \leq 2$ MBU characterization, $h + 1$ tree pebbling, contracts and complexity bounds for all nine core and four bootstrapped quantum data structures, $[8, 5, 3]_{\mathbb{F}_4}$ Hermitian curve AG Early-MBU DQI exactness, code-distance step-down, and convex-M^* distributed QRNS + GHZ cat-disentangling.

Appendix D indexes the exact Lean 4.35.0 declaration names machine-checked against this manuscript.

6.5 RQ5: Nine-Repository Open-Source Bug Audit and Standalone Linter

To test whether the hazards captured by $\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}$ occur in practice, we audited nine open-source quantum algorithm and compiler repositories—spanning elliptic-curve/RSA arithmetic [40, 48, 98], symmetric/hash circuits [89, 120], code-based ISD and DQI [34, 58, 60], and block-encoded QSVT libraries [33, 49]. As summarized in Table 7 (full repository-by-repository breakdown in Appendix D), LakefillQ's seven typing projections statically catch **14 latent correctness bugs**, including unguarded $x = 0$ Kaliski modular-inversion garbage (U1), missing PLU permutation uncomputation (U2), non-involutive ReVerC χ_5 uncomputation (U1), unguarded $d = 2$ QSVT walk sign flips when $U^2 \neq I$ (T-SubNorm), and uncanceled MBU phase kickbacks (U4). Our standalone Python linter (`ae/lakefillq_lint.py`) lifts these same checks onto external OpenQASM 3.0 and Qiskit circuits, catching 16/16 bug kernels (100% recall) with 0/14 false positives on verified clean circuits.

6.6 RQ6: Open-Problem Pareto Advances Across Three Execution Regimes

Table 8 highlights LakefillQ's end-to-end algorithmic results across **Single Machine**, **Parallel Machine** ($M^* \leq 256$), and **Serial Multishot** regimes (complete per-mode tables appear in Appendix C):

- **Public-Key Cryptanalysis (ECC-256, RSA-2048/4096):** Composing `PackedBezout` `ShiftZipper` (−512 Q), $s = 4$ `BinomialPebbleSchedule`, Select-SWAP, and totalized Hua squaring solves ECC-256 in **1,088 logical qubits** and 14.34×10^6 **Toffolis** (788 Q in ultra-compact mode), simultaneously beating prior low-width (2,124 Q, 198×10^6 Toffolis [48]) and low-Toffoli (2,871 Q [73]) designs. On RSA-2048/4096, `SkewBinary` `CarrySaveNumeral` achieves 1,364 Q / 2,644 Q (896 Q / 1,728 Q under Braun RNS-CRT; 1,924 Q single-focus) while preserving surface-code distance $d = 27$ instead of inflating to $d = 29$. In Parallel Machine mode, zero-carry QRNS + Law-U4 GHZ cat-disentangling cuts optical Bell pairs by 50.0%.
- **Symmetric, Post-Quantum, and DQI:** `DLog-5` and `BP32` S-boxes plus bidirectional half-cipher MBU Grover reach 261/325/389 Q (min-width) and 356 Q (min-gates AES-256); $h + 1$ sequential Merkle pebbling compresses SLH-DSA-128s to 4,320 Q (−34.8%); `StreamingSyndromeAccumulator` + QM4RI/DOOM compresses Classic McEliece-348864 to 7,744 Q (76.26× narrower) and BIKE-L1 to 36,969 Q (4,107× narrower); odd-half NTT

Table 8. Selected open-problem results compiled by LakefillQ vs. published state of the art (extended Single, Parallel, and Serial Multishot tables in Appendix C).

Problem / Workload	Published SOTA Baseline	LakefillQ Width	LakefillQ Toffolis / Metric	Key Gain
ECC-256 ECDLP	2,124 Q, 198M Toff. [48]	1,088 Q (788 min)	14.34M Toffolis	1.95× Q, 13.8× Toff.
RSA-2048 Factoring	2,052–4,098 Q ($d = 29$ seq.) [24, 40]	1,364 Q (896 RNS)	Preserves $d = 27$	344,304 BB-qLDPC Q
AES-256 Grover	512–1,376 Q [58, 70, 120]	356 Q (389 DLog-5)	32 TempAND/S-box	Bidirectional MBU
SLH-DSA-128s	6,624 Q simultaneous Merkle	4,320 Q	$h + 1$ two-pass zipper	−34.8% width
McEliece-348864 ISD	590,564 Q dense matrix [34]	7,744 Q	Streaming + QM4RI	76.26× narrower
Coherent OPI DQI	104,448 Q [60]	9,984 Q	Packed Bézout + pebbles	10.46× narrower
FeMoCo THC QSVT	2,046 PREP† TempANDs [69]	13 walk ancillas	91 PREP† (6,890/walk)	22.48× PREP†

compresses ML-KEM-768 / ML-DSA-65 to 3,072/5,888 Q; and coherent OPI DQI shrinks by 10.46× to 9,984 Q.

- **Quantum Chemistry (FeMoCo THC QSVT & Hubbard QPE):** QBlock Chebyshev reflection fusion (192 → 13 ancillas) and Law-U4 PREP† (2,046 → 91 TempANDs, 22.48×) lower FeMoCo walk cost to 6,890 TempANDs/step [69]. On noisy IBM Heron r2 Hubbard phase estimation [55, 71], Serial Multishot `reset_all` achieves 0.0795 Ha energy error and Parallel Machine achieves 0.0053 Ha, compared to 0.7540 Ha under monolithic Single Machine decoherence.

7 SCOPE BOUNDARIES, RELATED WORK, AND CONCLUSION

7.1 Negative Scope Boundaries: When Flat Imperative Circuits Prevail

Purely functional quantum data structures do not improve every quantum circuit; LakefillQ's cost model explicitly identifies three regimes where flat imperative circuits prevail:

(1) **High-Readout-Noise Hardware ($\epsilon_{\text{meas}} > 2\epsilon_{2Q}$) and Slow External Decoders ($\tau_{\text{dec}} > \tau_{\text{cycle}}$).** On devices with noisy dispersive readout or slow camera readout, mid-circuit MBU underperforms measurement-free 3-CZ Margolus RCCX or native RydbergCCZ; QPhysGrade automatically switches away from MBU in this regime (§5.1).

(2) **Small Static Lookup Tables ($M \leq 32$).** Below $M = 32$ entries, Batcher bitonic and crit-bit comparator overhead exceeds a flat Select-SWAP multiplexor.

(3) **QRAM-Free Asymptotic Query Lower Bounds.** Without speculative hardware QRAM, functional data structures reduce circuit width and constant-factor gate counts without violating query lower bounds. In particular, for ML-KEM and ML-DSA lattice sieving, `MonoidalFingerTreeFilter` cuts inner BKZ similarity checks to 255 Toffolis, yet QRAM-free Select-SWAP candidate scanning imposes $\Theta(N_{\text{vec}}^{1.5}) = 2^{0.31125\beta} > 2^{0.2920\beta}$ classical sieving [10, 23, 59]—honestly yielding a 0× asymptotic quantum sieving speedup without QRAM, just as $O(\kappa \log(1/\varepsilon))$ QSVT [42] and $\Theta(\sqrt{VM})$ minimum-spanning-tree bounds remain binding.

7.2 Related Work

Table 9 compares LakefillQ against representative quantum programming languages. Silq [17] pioneered automatic unitary uncomputation via syntactic `qfree`/`mfree` annotations, and Qrisp [102] automates unitary uncomputation via dependency graphs, but neither tracks bilinear X-basis MBU phase kickbacks (U4), domain-restricted algebraic inverses (U1), or hardware decoder/crosstalk grades (Φ_1–Φ_3). Twist [115] verifies entanglement purity (`pure`), and Tower [114] introduces pointer-based quantum radix trees with unitary pointer uncomputation; LakefillQ complements Tower with pointerless implicit trees, skew-binary carry-save numerals, packed

Table 9. Comparison of LakefillQ with prior quantum programming languages and verified compilers.

Language / System	Auto Uncomp.	M_X Phase (U4)	Canonical (U2)	HW Grade (Φ)	QBlock Fusion	Mechanized
Silq [17] / Qrisp [102]	Unitary	No	No	No	No	Paper / None
Twist [115]	Manual	No	No	No	No	Paper + Coq core
Tower [114]	Pointer unitary	No	Radix tries	No	No	Paper
Qunity [108]	Algebraic	No	No	No	Kraus only	Paper
QuRA / Proto-Quipper-R [27]	Affine width	No	No	Width only	No	Paper
VOQC [51]	Gate pass	No	No	Gate count	No	Coq
Guppy [66] / Qualtran [49]	Linear / Bloq	Dynamic / Est.	No	Bloq counts	Unfused	Rust / Python
Cobble [33]	Block-enc.	No	No	Subnorm. α	Unfused $d \cdot m$	Paper
LakefillQ (This Work)	**Unitary + MBU**	**XAG + Z3**	**Yes**	Φ_1–Φ_3 **+ ZX**	$m_{\text{anc}} + 1$ **+ U4**	**2,305 Lean 4**

Bézout zippers, and zero-Toffoli MBU phase tracking. Qunity [108] unifies control and Kraus semantics algebraically without substructural MBU or hardware grades, while Cobble [33] types continuous block-encodings without Law-U4 PREP† erasure or involutory Chebyshev ancilla fusion. Quipper [45], QWire [91], Proto-Quipper-R/QuRA [26, 27, 100], VOQC [51], Qualtran [49], and Guppy/HUGR [66] provide linear circuit description, affine width analysis, verified peephole optimization, or resource estimation, whereas LakefillQ integrates duoidal wire capabilities, XAG/FRAIG + Z3 verification, non-Clifford ZX e-graph extraction, and multi-QPU codegen in one mechanized calculus.

Within graded and modal type theory [1, 2, 8, 25, 37, 80, 86, 87, 111, 113], λ_{LFQ} combines uniqueness/linearity duality [77] and place-indexed distributed modalities [107] with two distinctly quantum phenomena absent from classical coeffect calculi: non-additive clean-ancilla reuse along chordal SSA graphs and nonlinear modal actions $r \triangleright e$ capturing Chebyshev phase drift and Zalka's $\sqrt{M}$ search penalty.

7.3 Conclusion

Purely functional data structures and fault-tolerant quantum circuits share a deep structural bond: both achieve efficiency by confining state changes to a small focus or logarithmic spine rather than mutating aliased memory in place. By uniting duoidal wire capabilities ($0 < \mathsf{BorrowZ} < \mathsf{TaintedX} < \mathsf{Lin1} < \mathsf{Uniq1}$) and physical hardware grades in a single compositional judgment $\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}$, LakefillQ enables annotation-free verification of MBU phase cancellation, non-Clifford ZX e-graph lowering, and Pareto-optimal circuits across cryptography, algebraic-geometry DQI, post-quantum security, and quantum chemistry.

REFERENCES

[1] Andreas Abel and Jean-Philippe Bernardy. 2020. A Unified View of Modalities in Type Systems. *Proceedings of the ACM on Programming Languages* 4, ICFP (2020), 90:1–90:28.

[2] Andreas Abel, Nils Anders Danielsson, and Oskar Eriksson. 2023. A Graded Modal Dependent Type Theory with a Universe and Erasure, Formalized. *Proceedings of the ACM on Programming Languages* 7, ICFP (2023), 220:920–220:954.

[3] Thorsten Altenkirch and Jonathan Grattage. 2005. A Functional Quantum Programming Language (QML). In *Proceedings of the 20th Annual IEEE Symposium on Logic in Computer Science (LICS)*. 249–258.

[4] Andris Ambainis. 2007. Quantum Walk Algorithm for Element Distinctness. *SIAM J. Comput.* 37, 1 (2007), 210–239.

[5] Andris Ambainis. 2012. Variable Time Amplitude Amplification and a Faster Quantum Algorithm for Solving Systems of Linear Equations. In *29th International Symposium on Theoretical Aspects of Computer Science (STACS) (LIPIcs, Vol. 14)*. 636–647.

[6] Matthew Amy, Dmitri Maslov, and Michele Mosca. 2014. Polynomial-Time T-Depth Optimization of Clifford+T Circuits via Matroid Partitioning. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems* 33, 10 (2014), 1476–1489.

[7] Matthew Amy, Olivia Di Matteo, Vlad Gheorghiu, Michele Mosca, Alex Parent, and John Schanck. 2016. Estimating the Cost of Generic Quantum Pre-Image Attacks on SHA-2 and SHA-3. In *Selected Areas in Cryptography (SAC) (LNCS, Vol. 10532)*. 317–337.

[8] Robert Atkey. 2018. Syntax and Semantics of Quantitative Type Theory (QTT). In *Proceedings of the 33rd Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*. 56–65.

[9] Ryan Babbush, Craig Gidney, Dominic W. Berry, Nathan Wiebe, Jarrod McClean, Alexandru Paler, Austin Fowler, and Hartmut Neven. 2018. Encoding Electronic Spectra in Quantum Circuits with Linear T Complexity. *Physical Review X* 8, 4 (2018), 041015.

[10] Anja Becker, Léo Ducas, Nicolas Gama, and Thijs Laarhoven. 2016. New Directions in Nearest Neighbor Searching with Applications to Lattice Sieving. In *Proceedings of the 27th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*. 10–24.

[11] Charles H. Bennett. 1973. Logical Reversibility of Computation. *IBM Journal of Research and Development* 17, 6 (1973), 525–532.

[12] Charles H. Bennett. 1989. Time/Space Trade-Offs for Reversible Computation. *SIAM J. Comput.* 18, 4 (1989), 766–776.

[13] Jon Louis Bentley and James B. Saxe. 1980. Decomposable Searching Problems I: Static-to-Dynamic Transformation. *Journal of Algorithms* 1, 4 (1980), 301–358.

[14] Daniel J. Bernstein. 2009. Cost Analysis of Hash Collisions: Will Quantum Computers Make SHARCS Obsolete?. In *SHARCS 2009: Workshop Record*. 105–116.

[15] Daniel J. Bernstein and Bo-Yin Yang. 2019. Fast Constant-Time GCD Computation and Modular Inversion. *IACR Transactions on Cryptographic Hardware and Embedded Systems (TCHES)* 2019, 3 (2019), 340–398.

[16] Dominic W. Berry, Craig Gidney, Mario Motta, Jarrod R. McClean, and Ryan Babbush. 2019. Qubitization of Arbitrary Basis Quantum Chemistry Leveraging Sparsity and Low Rank Factorization. *Quantum* 3 (2019), 208.

[17] Benjamin Bichsel, Maximilian Baader, Timon Gehr, and Martin Vechev. 2020. Silq: A High-Level Quantum Language with Safe Uncomputation and Intuitive Semantics. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. 286–300.

[18] Dolev Bluvstein, Simon J. Evered, Alexandra A. Geim, Sophie H. Li, Hengyun Zhou, Tom Manovitz, Sepehr Ebadi, Madelyn Cain, Marcin Kalinowski, Dominik Hangleiter, J. Pablo Bonilla Ataides, Nishad Maskara, Iris Cong, Xun Gao, Pedro Sales Rodriguez, Thomas Karolyshyn, Giulia Semeghini, Michael J. Gullans, Markus Greiner, Vladan Vuletić, and Mikhail D. Lukin. 2024. Logical Quantum Processor Based on Reconfigurable Atom Arrays. *Nature* 626 (2024), 58–65.

[19] Dolev Bluvstein, Pedro Sales Rodriguez, and Mikhail D. Lukin. 2026. Zoned Fault-Tolerant Architecture with Coherent Neutral Atoms. *Nature* 649 (2026). With native 3-qubit Rydberg blockade CCZ gates, arXiv:2609.11159 (2026).

[20] Joan Boyar and René Peralta. 2012. A Small Depth-16 Circuit for the AES S-Box. In *IFIP International Information Security and Privacy Conference (SEC)*. Springer, 287–298.

[21] W. Braun and M. Rem. 1983. *A Logarithmic Implementation of Flexible Arrays.* Technical Report MR83/4. Eindhoven University of Technology.

[22] Sergey Bravyi, Andrew W. Cross, Jay M. Gambetta, Dmitri Maslov, Patrick Rall, and Theodore J. Yoder. 2024. High-Threshold and Low-Overhead Fault-Tolerant Quantum Memory. *Nature* 627 (2024), 778–782.

[23] André Chailloux and Johanna Loyer. 2021. Lattice Sieving via Quantum Random Walks. In *Advances in Cryptology – ASIACRYPT 2021 (LNCS, Vol. 13093)*. 63–91.

[24] Clémence Chevignard, Pierre-Alain Fouque, and André Schrottenloher. 2024. Reducing the Number of Qubits in Quantum Factoring. In *Advances in Cryptology – ASIACRYPT 2024 (LNCS, Vol. 15489)*. 3–35.

[25] Pritam Choudhury, Harley Eades III, Richard A. Eisenberg, and Stephanie Weirich. 2021. A Graded Dependent Type System with a Usage-Aware Semantics. *Proceedings of the ACM on Programming Languages* 5, POPL (2021), 50:1–50:32.

[26] Andrea Colledan and Ugo Dal Lago. 2024. Circuit Width Estimation via Effect Typing and Linear Dependency (QuRA). In *Proceedings of the 33rd European Symposium on Programming (ESOP) (LNCS, Vol. 14577)*. 3–30. Extended with Proto-Quipper-R in POPL 2025.

[27] Andrea Colledan and Ugo Dal Lago. 2025. On Dependent Types and Quantum Resource Analysis. *Proceedings of the ACM on Programming Languages* 9, POPL (2025), 14:387–14:417.

[28] Steven A. Cuccaro, Thomas G. Draper, Samuel A. Kutin, and David Petrie Moulton. 2004. A New Quantum Ripple-Carry Addition Circuit. arXiv:quant-ph/0410184.

[29] Ryan Culpepper and Matthias Felleisen. 2012. Fortifying Macros. *Journal of Functional Programming* 22, 4–5 (2012), 439–476.

[30] David Deutsch, Artur Ekert, Richard Jozsa, Chiara Macchiavello, Sandu Popescu, and Anna Sanpera. 1996. Quantum Privacy Amplification and the Security of Quantum Cryptography over Noisy Channels (DEJMPS). *Physical Review*

Letters 77, 13 (1996), 2818–2821.

[31] Ross Duncan, Aleks Kissinger, Simon Perdrix, and John van de Wetering. 2020. Graph-Theoretic Simplification of Quantum Circuits with the ZX-Calculus. *Quantum* 4 (2020), 279.

[32] Martin Ekerå and Johan Håstad. 2017. Quantum Algorithms for Computing Short Discrete Logarithms and Factoring RSA Integers. In *Post-Quantum Cryptography (PQCrypto 2017) (LNCS, Vol. 10346)*. 347–363.

[33] ELLORA and Cobble Authors. 2026. Cobble: Compositional Block-Encoding Language and Polynomial/Sum Fusion Compiler. In *Proceedings of the 47th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. arXiv:2511.01736.

[34] Andre Esser, Sergi Ramos-Calderer, Emanuele Bellini, José Ignacio Latorre, and Marc Manzano. 2022. An Optimized Quantum Implementation of ISD on Scalable Quantum Resources. In *Post-Quantum Cryptography (PQCrypto 2022) (LNCS, Vol. 13512)*. 262–288.

[35] Gui-Liang Feng and T. R. N. Rao. 1993. Decoding Algebraic-Geometric Codes up to the Designed Minimum Distance. *IEEE Transactions on Information Theory* 39, 1 (1993), 37–45.

[36] Peng Fu, Kohei Kishida, Neil J. Ross, and Peter Selinger. 2020. A Tutorial Introduction to Quantum Circuit Programming in Dependently Typed Proto-Quipper (Proto-Quipper-D). In *Proceedings of the 12th International Conference on Reversible Computation (RC) (LNCS, Vol. 12227)*. 153–168.

[37] Marco Gaboardi, Shin ya Katsumata, Dominic Orchard, Flavien Breuvart, and Tarmo Uustalu. 2016. Combining Effects and Coeffects via Grading. In *Proceedings of the 21st ACM SIGPLAN International Conference on Functional Programming (ICFP 2016)*. ACM, New York, NY, USA, 476–489.

[38] Craig Gidney. 2018. Halving the Cost of Quantum Addition. *Quantum* 2 (2018), 74.

[39] Craig Gidney. 2025. How to Factor 2048 Bit RSA Integers with Less Than a Million Noisy Qubits. arXiv:2505.15917.

[40] Craig Gidney and Martin Ekerå. 2021. How to Factor 2048 Bit RSA Integers in 8 Hours Using 20 Million Noisy Qubits. *Quantum* 5 (2021), 433.

[41] Craig Gidney, Noah Shutty, and Cody Jones. 2024. Magic State Cultivation: Growing T States as Cheap as CNOT Gates. arXiv:2409.17595.

[42] András Gilyén, Yuan Su, Guang Hao Low, and Nathan Wiebe. 2019. Quantum Singular Value Transformation and Beyond: Exponential Improvements for Quantum Matrix Arithmetics. In *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing (STOC)*. 193–204.

[43] Google Quantum AI. 2025. Quantum Error Correction Below the Surface Code Threshold (Willow). *Nature* 638 (2025), 920–926. With real-time closed-loop decoding in arXiv:2605.04892 (2026).

[44] Élie Gouzien, Diego Ruiz, Flavien Le Régent, Jérémie Guillaud, and Nicolas Sangouard. 2023. Performance Analysis of a Repetition Cat Code Architecture: Computing 256-Bit Elliptic Curve Logarithm in 9 Hours with 126,133 Cat Qubits. *Physical Review Letters* 131, 4 (2023), 040602.

[45] Alexander S. Green, Peter LeFanu Lumsdaine, Neil J. Ross, Peter Selinger, and Benoît Valiron. 2013. Quipper: A Scalable Quantum Programming Language. In *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. 333–342.

[46] Sebastian Hack, Daniel Grund, and Gerhard Goos. 2006. Register Allocation for Programs in SSA-Form. In *Proceedings of the 15th International Conference on Compiler Construction (CC) (LNCS, Vol. 3923)*. Springer, 247–262.

[47] Thomas Häner, Samuel Jaques, Norbert M. Linke, and IonQ Fault-Tolerant Architecture Team. 2026. Computing 256-Bit Elliptic Curve Discrete Logarithms in 26 Days on a Fault-Tolerant Trapped-Ion Quantum Computer with 20,000 Qubits. arXiv:2609.05625, Q102 Walking-Cat qLDPC blueprint (1,457 logical qubits, 39M CCZ, 19,397 physical ions).

[48] Thomas Häner, Samuel Jaques, Michael Naehrig, Martin Roetteler, and Mathias Soeken. 2020. Improved Quantum Circuits for Elliptic Curve Discrete Logarithms. In *Post-Quantum Cryptography (PQCrypto 2020) (LNCS, Vol. 12100)*. 425–444.

[49] Matthew P. Harrigan, Tanuj Khattar, Charles Yuan, Anurudh Peduri, Noureldin Yosri, Fionn D. Malone, Ryan Babbush, and Nicholas C. Rubin. 2024. Expressing and Analyzing Quantum Algorithms with Qualtran. arXiv:2409.04643.

[50] Bastiaan Heeren, Jurriaan Hage, and S. Doaitse Swierstra. 2003. Scripting the Type Inference Process. In *Proceedings of the Eighth ACM SIGPLAN International Conference on Functional Programming (ICFP '03)*. ACM, New York, NY, USA, 3–13.

[51] Kesha Hietala, Robert Rand, Shih-Han Hung, Xiaodi Wu, and Michael Hicks. 2021. A Verified Optimizer for Quantum Circuits. *Proceedings of the ACM on Programming Languages* 5, POPL (2021), 1–29.

[52] Ralf Hinze and Ross Paterson. 2006. Finger Trees: A Simple General-Purpose Data Structure. *Journal of Functional Programming* 16, 2 (2006), 197–217.

[53] Robert Hood and Robert Melville. 1981. Real-Time Queue Operations in Pure LISP. *Inform. Process. Lett.* 13, 2 (1981), 50–54.

[54] Loo-Keng Hua. 1949. Some Properties of a Sfield. *Proc. Nat. Acad. Sci. U. S. A.* 35, 9 (1949), 533–537.

[55] Zhenyu Huang and Siwei Sun. 2022. Synthesizing Quantum Circuits of AES with Lower T-Depth and Less Qubits. In *Advances in Cryptology – ASIACRYPT 2022 (LNCS, Vol. 13793)*. 614–644.

[56] Gérard Huet. 1997. The Zipper. *Journal of Functional Programming* 7, 5 (1997), 549–554.

[57] IBM Quantum. 2026. Dynamic Circuits and Heavy-Hex Fault-Tolerant Primitives on Heron r2. *Physical Review* (2026). https://doi.org/10.1103/jcdk-q3xc

[58] Samuel Jaques, Michael Naehrig, Martin Roetteler, and Fernando Virdia. 2020. Implementing Grover Oracles for Quantum Key Search on AES and LowMC. In *Advances in Cryptology – EUROCRYPT 2020 (LNCS, Vol. 12106)*. 280–310.

[59] Samuel Jaques and John M. Schanck. 2019. Quantum Cryptanalysis in the RAM Model: Claw-Finding Attacks on SIKE. In *Advances in Cryptology – CRYPTO 2019 (LNCS, Vol. 11692)*. 32–61.

[60] Stephen P. Jordan, Noah Shutty, Mary Wootters, Adam Zalcman, Alexander Schmidhuber, Robbie King, Sergei V. Isakov, and Ryan Babbush. 2025. Optimization by Decoded Quantum Interferometry. *Nature* (2025). arXiv:2408.08292.

[61] Burton S. Kaliski Jr. 1995. The Montgomery Inverse and Its Applications. *IEEE Trans. Comput.* 44, 8 (1995), 1064–1065.

[62] Shinichi Kawamura, Masanobu Koike, Fumihiko Sano, and Atsushi Shimbo. 2000. Cox–Rower Architecture for Fast Parallel Montgomery Multiplication. In *Advances in Cryptology – EUROCRYPT 2000 (LNCS, Vol. 1807)*. 523–538.

[63] Andrew W. Keep and R. Kent Dybvig. 2013. A Nanopass Framework for Commercial Compiler Development. In *Proceedings of the 18th ACM SIGPLAN International Conference on Functional Programming (ICFP '13)*. ACM, New York, NY, USA, 343–350.

[64] Tanuj Khattar and Ryan Babbush. 2025. *Fault-Tolerant Resource Estimates and Syndrome-Measurement Optimizations for Decoded Quantum Interferometry*. Technical Report. Google Quantum AI.

[65] Aleks Kissinger and John van de Wetering. 2020. Reducing the Number of Non-Clifford Gates in Quantum Circuits. *Physical Review A* 102, 2 (2020), 022406.

[66] Mark Koch, Agustin Borgna, Seyon Sivarajah, Alan Lawrence, Alec Edgington, Douglas Wilson, and Ross Duncan. 2025. HUGR, Guppy, and TKET2: A Hierarchical Unified Graph Intermediate Representation for Dynamic Quantum-Classical Programs. arXiv:2510.13082.

[67] Richard Královič. 2004. Time and Space Complexity of Reversible Pebbling. *RAIRO - Theoretical Informatics and Applications* 38, 2 (2004), 137–161.

[68] L. Kraniotis and T. Theoharis. 2003. Note on the Reversible Pebbling of Complete Binary Trees. *Inform. Process. Lett.* 88 (2003), 239–242.

[69] Joonho Lee, Dominic W. Berry, Craig Gidney, William J. Huggins, Jarrod R. McClean, Nathan Wiebe, and Ryan Babbush. 2021. Even More Efficient Quantum Computations of Chemistry Through Tensor Hypercontraction. *PRX Quantum* 2, 3 (2021), 030305.

[70] Da Lin, Zejun Xiang, Xiangyong Zeng, and Shasha Zhang. 2025. Quantum Circuit Implementations of AES with Minimal Qubits via In-Place S-Boxes. arXiv:2503.06097 (264/328/392-qubit AES-128/192/256 baseline).

[71] Lin Lin and Yu Tong. 2022. Heisenberg-Limited Ground-State Energy Estimation for Early Fault-Tolerant Quantum Computers. *PRX Quantum* 3, 1 (2022), 010318.

[72] Daniel Litinski. 2019. A Game of Surface Codes: Large-Scale Quantum Computing with Lattice Surgery. *Quantum* 3 (2019), 128.

[73] Daniel Litinski. 2023. How to Compute a 256-Bit Elliptic Curve Private Key with Only 50 Million Toffoli Gates. arXiv:2306.08585.

[74] Guang Hao Low and Isaac L. Chuang. 2019. Hamiltonian Simulation by Qubitization. *Quantum* 3 (2019), 163.

[75] Guang Hao Low, Vadym Kliuchnikov, and Luke Schaeffer. 2024. Trading T-Gates for Dirty Qubits in State Preparation and Unitary Synthesis. *Quantum* 8 (2024), 1375. arXiv:1812.00954.

[76] Guillaume Marceau, Kathi Fisler, and Shriram Krishnamurthi. 2011. Measuring the Effectiveness of Error Messages Designed for Novice Programmers. In *Proceedings of the 42nd ACM Technical Symposium on Computer Science Education (SIGCSE '11)*. ACM, New York, NY, USA, 499–504.

[77] Daniel Marshall, Michael Vollmer, and Dominic Orchard. 2022. Linearity and Uniqueness: An Entente Cordiale. In *Programming Languages and Systems (ESOP 2022) (LNCS, Vol. 13240)*. Springer, Cham, 346–375.

[78] Dmitri Maslov. 2016. Advantages of Using Relative-Phase Toffoli Gates with an Application to Multiple Control Toffoli Optimization (RCCX). *Physical Review A* 93, 2 (2016), 022311.

[79] Conor McBride. 2001. The Derivative of a Regular Type Is Its Type of One-Hole Contexts. (2001). Unpublished manuscript.

[80] Benjamin Moon, Harley Eades III, and Dominic Orchard. 2021. Graded Modal Dependent Type Theory. In *Programming Languages and Systems (ESOP 2021) (LNCS, Vol. 12648)*. Springer, Cham, 455–484.

[81] Donald R. Morrison. 1968. PATRICIA—Practical Algorithm To Retrieve Information Coded in Alphanumeric. *J. ACM* 15, 4 (1968), 514–534.

[82] Eugene W. Myers. 1983. An Applicative Random-Access Stack. *Inform. Process. Lett.* 17, 5 (1983), 241–248.

[83] Chris Okasaki. 1997. Three Algorithms on Braun Trees. *Journal of Functional Programming* 7, 6 (1997), 661–666.

[84] Chris Okasaki. 1998. *Purely Functional Data Structures.* Cambridge University Press.

[85] Chris Okasaki and Andy Gill. 1998. Fast Mergeable Integer Maps. In *Workshop on ML*. 77–86.

[86] Dominic Orchard, Vilem-Benjamin Liepelt, and Harley Eades III. 2019. Quantitative Program Reasoning with Graded Modal Types. *Proceedings of the ACM on Programming Languages* 3, ICFP (2019), 110:1–110:30.

[87] Dominic Orchard, Philip Wadler, and Harley Eades III. 2020. Unifying Graded and Parameterised Monads. In *Proceedings of the 8th Workshop on Mathematically Structured Functional Programming (MSFP 2020) (EPTCS, Vol. 317)*. Open Publishing Association, Waterloo, Australia, 18–38.

[88] Anouk Paradis, Benjamin Bichsel, Sumit Gulwani, and Martin Vechev. 2021. Unqomp: Synthesizing Uncomputation in Quantum Circuits. In *Proceedings of the 42nd ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. 222–236.

[89] Alex Parent, Martin Roetteler, and Krysta M. Svore. 2017. ReVerC: A Reversible Circuit Compiler for Space-Constrained Quantum Computations. In *Reversible Computation (RC 2017) (LNCS, Vol. 10301)*. 3–20.

[90] Zvonimir Pavlinović, Tim King, and Thomas Wies. 2014. Finding Minimum Type Error Sources. In *Proceedings of the 2014 ACM International Conference on Object Oriented Programming Systems Languages & Applications (OOPSLA '14)*. ACM, New York, NY, USA, 525–542.

[91] Jennifer Paykin, Robert Rand, and Steve Zdancewic. 2017. QWIRE: A Core Language for Quantum Circuits. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages (POPL)*. 846–858.

[92] Simone Perriello, Alessandro Barenghi, and Gerardo Pelosi. 2023. Improving Quantum Information Set Decoding. *ACM Transactions on Quantum Computing* 4, 4 (2023), 1–40. Extended with Lee–Brickell quantum circuits in ACNS 2024.

[93] Massimiliano Poletto and Vivek Sarkar. 1999. Linear Scan Register Allocation. *ACM Transactions on Programming Languages and Systems* 21, 5 (1999), 895–913.

[94] Justin Pombrio and Shriram Krishnamurthi. 2018. Inferring Type Rules for Syntactic Sugar. *Proceedings of the ACM on Programming Languages* 2, PLDI (2018), 81:1–81:28.

[95] Oded Regev. 2025. An Efficient Quantum Factoring Algorithm. *J. ACM* 72, 1 (2025), 1–13. arXiv:2308.06572.

[96] Joost Renes, Craig Costello, and Lejla Batina. 2016. Complete Addition Formulas for Prime Order Elliptic Curves. In *Advances in Cryptology – EUROCRYPT 2016 (LNCS, Vol. 9665)*. 403–428.

[97] Francisco Rios and Peter Selinger. 2018. A Categorical Model for a Quantum Circuit Description Language (Proto-Quipper-M). In *Proceedings of the 14th International Conference on Quantum Physics and Logic (QPL) (EPTCS, Vol. 266)*. 164–178.

[98] Martin Roetteler, Michael Naehrig, Krysta M. Svore, and Kristin Lauter. 2017. Quantum Resource Estimates for Computing Elliptic Curve Discrete Logarithms. In *Advances in Cryptology – ASIACRYPT 2017 (LNCS, Vol. 10625)*. 241–270.

[99] Neil J. Ross and Peter Selinger. 2016. Optimal Ancilla-Free Clifford+T Approximation of z-Rotations (`gridsynth`). *Quantum Information and Computation* 16, 11–12 (2016), 901–953.

[100] Ken Sakayori, Andrea Colledan, and Ugo Dal Lago. 2025. Proto-Quipper-C: Closure-Indexed Resource Analysis for Higher-Order Quantum Circuit Programs. arXiv:2511.22419.

[101] Eric L. Seidel, Ranjit Jhala, and Westley Weimer. 2016. Dynamic Witnesses for Static Type Errors (or, Ill-Typed Programs Usually Go Wrong). In *Proceedings of the 21st ACM SIGPLAN International Conference on Functional Programming (ICFP '16)*. ACM, New York, NY, USA, 228–242.

[102] Raphael Seidel, Sebastian Bock, René Zander, Matic Petrič, Niklas Steinmann, Nikolay Tcholtchev, and Manfred Hauswirth. 2024. Qrisp: A Framework for Compilable High-Level Programming of Gate-Based Quantum Computers. arXiv:2406.14792.

[103] Peter Selinger and Benoît Valiron. 2006. A Lambda Calculus for Quantum Computation with Classical Control. *Mathematical Structures in Computer Science* 16, 3 (2006), 527–552.

[104] Ravi Sethi and Jeffrey D. Ullman. 1970. The Generation of Optimal Code for Arithmetic Expressions. *J. ACM* 17, 4 (1970), 715–728.

[105] Henning Stichtenoth. 2009. *Algebraic Function Fields and Codes* (2nd ed.). Graduate Texts in Mathematics, Vol. 254. Springer, Berlin, Germany.

[106] Barbara M. Terhal. 2015. Quantum Error Correction for Quantum Memories. *Reviews of Modern Physics* 87, 2 (2015), 307–346.

[107] Tom Murphy VII, Karl Crary, and Robert Harper. 2008. Type-Safe Distributed Programming with ML5. In *Trustworthy Global Computing (TGC 2007) (LNCS, Vol. 4912)*. Springer, Berlin, Heidelberg, 108–123.

[108] Finn Voichick, Liyi Li, Robert Rand, and Michael Hicks. 2023. Qunity: A Unified Language for Quantum and Classical Computing. *Proceedings of the ACM on Programming Languages* 7, POPL (2023), 1–31. Extended with Kraus-to-unitary compilation in PACMPL 9(OOPSLA), 2025.
[109] Michael D. Vose. 1991. A Linear Algorithm for Generating Random Numbers with a Given Distribution. *IEEE Transactions on Software Engineering* 17, 9 (1991), 972–975.
[110] Alastair J. Walker. 1977. An Efficient Method for Generating Discrete Random Variables with General Distributions. *ACM Trans. Math. Software* 3, 3 (1977), 253–256.
[111] James Wood and Robert Atkey. 2022. A Framework for Substructural Type Systems. In *Programming Languages and Systems (ESOP 2022) (LNCS, Vol. 13240)*. Springer, Cham, 376–402.
[112] Zeju Xiang, Xiangyong Zeng, Dongdai Lin, Zhenzhen Bao, and Shasha Zhang. 2020. Optimizing Implementations of Linear Layers. *IACR Transactions on Symmetric Cryptology (ToSC)* 2020, 2 (2020), 120–145.
[113] Shin ya Katsumata. 2014. Parametric Effect Monads and Semantics of Effect Systems. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL 2014)*. ACM, New York, NY, USA, 633–645.
[114] Charles Yuan and Michael Carbin. 2022. Tower: Data Structures in Quantum Superposition. *Proceedings of the ACM on Programming Languages* 6, OOPSLA2 (2022), 259–288.
[115] Charles Yuan, Christopher McNally, and Michael Carbin. 2022. Twist: Sound Reasoning for Purity and Entanglement in Quantum Programs. *Proceedings of the ACM on Programming Languages* 6, POPL (2022), 1–32.
[116] Christof Zalka. 1999. Grover's Quantum Searching Algorithm Is Optimal. *Physical Review A* 60, 4 (1999), 2746–2751.
[117] Mark Zhandry. 2019. How to Record Quantum Queries, and Applications to Quantum Indifferentiability. In *Advances in Cryptology – CRYPTO 2019 (LNCS, Vol. 11693)*. 239–268.
[118] Danfeng Zhang and Andrew C. Myers. 2014. Toward General Diagnosis of Static Errors. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '14)*. ACM, New York, NY, USA, 569–581.
[119] Junyu Zhou and Collaborators. 2026. DShor: Distributed Quantum Factoring via Approximate Residue Number Systems and GHZ Lookup Tables. arXiv:2605.03951 (6-module distributed RSA-2048 factoring).
[120] Jian Zou, Zihao Wei, Siwei Sun, Ximeng Liu, and Wenling Wu. 2020. Quantum Circuit Implementations of AES with Fewer Qubits. In *Advances in Cryptology – ASIACRYPT 2020 (LNCS, Vol. 12492)*. 697–726.

A HOST-LANGUAGE REALIZATIONS OF QUANTUM DATA STRUCTURES

While `#lang lakefillq` serves as our primary functional surface language in the main body, researchers integrating low-width quantum circuits into existing classical toolchains often embed reversible circuit combinators in typed host languages [3, 36, 97, 103]. Following McBride's derivative view of one-hole contexts [79] and Okasaki–Gill Braun indexing [85], our artifact provides verified host embeddings in Rust (`no_std` gate-level emitter benchmarking Cuccaro [28] vs. Gidney [38] adders), Haskell (`-XLinearTypes`), OCaml functors, and Lean 4.35.0. Listings 7 and 8 illustrate the gate-level Rust MBU ripple-carry primitive and the linear Haskell/OCaml/Lean type signatures corresponding to $\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}$.

Listing 7. Gate-level Rust `no_std` emitter for Gidney's degree-2 MBU ripple-carry adder ($W - 1$ forward Toffolis, 0 reverse Toffolis).

```rust
pub fn gidney_ripple_add_in_place_mbu<const W: usize>(
    a: &mut QReg<W>, b: &mut QReg<W>, carry: &mut QReg<W>,
) {
    CCX(a[0], b[0], &mut carry[0]);
    for i in 1..(W - 1) {
        CNOT(carry[i - 1], &mut a[i]);
        CNOT(carry[i - 1], &mut b[i]);
        CCX(a[i], b[i], &mut carry[i]);
        CNOT(carry[i - 1], &mut carry[i]);
    }
    CNOT(carry[W - 2], &mut b[W - 1]);
    CNOT(a[W - 1], &mut b[W - 1]);
    for i in (1..(W - 1)).rev() {
        CNOT(carry[i - 1], &mut carry[i]);
        if MEAS_X_RESET(&mut carry[i]) { CZ(a[i], b[i]); }
        CNOT(carry[i - 1], &mut a[i]);
    }
```

```
    if MEAS_X_RESET(&mut carry[0]) { CZ(a[0], b[0]); }
}
```

Listing 8. Linear Haskell (`-XLinearTypes`) GADT signature encoding the five-point wire capability scale and scoped focus-zipper invariant.

```
{-# LANGUAGE DataKinds, GADTs, LinearTypes #-}
data Cap = Cap0 | BorrowZ | TaintedX | Lin1 | Uniq1
data QWire (k :: Cap) where MkWire :: Int -> QWire k

withFocusZipper :: (QVec w 'Lin1 %1 -> QVec c 'BorrowZ %1 -> (QVec w 'Lin1, QVec c 'BorrowZ))
                -> QVec w 'Lin1 %1 -> QVec c 'Lin1 %1 -> (QVec w 'Lin1, QVec c 'Lin1)
```

B PROOFS OF FORMAL PROPOSITIONS, LEMMAS, AND THEOREMS

We summarize the mathematical proofs of all formal statements in the main paper; every proof is machine-checked without axioms in Lean 4.35.0 (`proofs/*.lean`, indexed in Appendix D).

Proof of Proposition 4.1 (Degree-2 Boundary of Clifford MBU).. Applying H and Z-basis measurement $m \in \{0,1\}$ to $|f(x)\rangle$ maps $\sum_x \alpha_x |x\rangle|f(x)\rangle \mapsto \frac{1}{\sqrt{2}} \sum_x \alpha_x (-1)^{m\cdot f(x)} |x\rangle|m\rangle$. When $m = 1$, canceling $(-1)^{f(x)}$ requires a diagonal unitary $D = \mathrm{diag}((-1)^{g(x)})$. By the Clifford hierarchy characterization of diagonal gates [7], a diagonal operator over n qubits is generated by $\{\mathrm{Z}, \mathrm{S}, \mathrm{CZ}\}$ if and only if its phase exponent has polynomial degree ≤ 2 over $\mathbb{Z}_4$, which for Boolean $\{\pm 1\}$ phases $(-1)^{g(x)}$ holds iff $\deg_{\mathbb{F}_2}(g) \leq 2$. When $\deg_{\mathbb{F}_2}(f) \geq 3$ (minimally $f(x) = x_1x_2x_3$ on $n = 3$), the minimum Hamming distance from any cubic Boolean function to the second-order Reed–Muller code $\mathcal{RM}(2,3)$ is $2^{3-3} = 1$ out of 8 truth-table entries, bounding the inner-product fidelity $|\frac{1}{8}\sum_{x\in\{0,1\}^3}(-1)^{f(x)\oplus g(x)}|^2 \leq (6/8)^2 = 9/16 = 0.5625$. □

Proof of Proposition 4.2 (Sequential Two-Pass Tree Pebbling). Let $R(h)$ denote the minimum W-qubit digest registers to compute the root of a complete height-h Merkle tree. Simultaneous evaluation computes left child ($R(h-1)$), holds its W-qubit output while computing right child, and hashes into a parent register, requiring $2h+1$ registers across simultaneous active spines. Factoring the parent sponge/hash absorption into S_{left} followed by S_{right} computes the left child in $R(h-1)$ registers, XOR/absorbs it into the persistent parent state, and uncomputes the left child completely back to $|0\rangle$ before evaluating the right child in the *same* $R(h-1)$ pool, yielding $R(h) = \max(R(h-1)+1, R(h-1)) = h+1$ [68, 104]. □

Proofs of Propositions 4.3–4.11 (Quantum Data Structures). (4.3) Rule T-FocusZipper restricts mutable Lin1 capabilities to the w-bit focus, so zero Toffolis touch the $(C-1)w$ BorrowZ wires; constant shifts permute compile-time wire indices (0 gates), and planar layout places cold limbs in a $\sqrt{Q_{\text{cold}}} \times \sqrt{Q_{\text{cold}}}$ memory block bordered by $2\lceil\sqrt{Q_{\text{cold}}}\rceil$ routing tiles. (4.4) Writing $X = \sum_{k=0}^{C-1}(L_k + d_k 2^w)2^{kw}$ with $L_k \in [0, 2^w - 1]$ and $d_k \in \{0,1,2\}$ preserves integer value modulo 2^{Cw} under window carry folding $d'_k = d_k + c_{\text{out}} \leq 2$ (maintained via $O(1)$-depth inter-window local carry fold), and a single terminal ripple sweep sets all $d_k \mapsto 0$ bijectively. (4.5) By the Braun invariant $0 \leq |L| - |R| \leq 1$, even/odd index parity fixes subtree routing statically without pointer qubits; associative binary folds execute in $\lceil \log_2 k \rceil$ rounds using $(k-1)w$ forward TempANDs. (4.6) Throughout Kaliski/Bernstein–Yang binary GCD [15, 61], $|u|\cdot|g| + |v|\cdot|f| = p$, ensuring $\mathrm{len}(u) + \mathrm{len}(f) \leq N+1$; packing u MSB-down and f LSB-up in one N-bit register (plus focus guard bits) avoids overlap, half-domain cofactor folding recovers $a^{-1} \bmod p$ bijectively on $\mathbb{F}_p^*$ (checked across all 65,520 slopes of $\mathbb{F}_{65521}^*$), and indicator corrections on $x \in \{0,1\}$ make Hua squaring total on $\mathbb{F}_p$. (4.7) Standard induction on (s,d) [12, 67] gives $L(s,d) = L(s,d-1) + L(s-1,d) = \binom{s+d}{s}$; at $s = 4, d = 4$, $\binom{8}{4} = 70 \geq 32$ Jump-16 windows in 64 qubits and 109 window evaluations. (4.8–4.11) Monoidal

prefix associativity validates 255-Toffoli early rejection in Two-Level VTAA; row-pivot tracking in $r\lceil \log_2 r \rceil$ bits makes Bruhat LU total across all 20, 160 matrices in $\mathrm{GL}_4(\mathbb{F}_2)$ while column streaming avoids r^2 allocation [92]; Morrison crit-bit splits depend only on highest differing key bits, yielding permutation-invariant quantum states [117] with 14, 080-Toffoli amortized parallel-walker bitonic batches (14, 417, 920 total vs. 23, 855, 104 Bentley–Saxe online Toffolis); and Walker–Vose bucket weights equal p_i after 1-comparator swap on $2w + 2\mu + 1$ qubits. □

Proofs of Theorems 3.1–3.4, Lemmas 5.1–5.2, Proposition 5.5, and Corollary 5.4. (3.1) Structural induction on all 20 surface constructors eliminates every macro into $\lambda^{\mathrm{kern}}_{\mathrm{LFQ}}$ (5/5 positive reference `.lfq` programs plus Hermitian AG DQI compile cleanly, 7/7 negative violation programs are rejected, 3/3 hardware-violating programs are restored by `repairPhysGrade`, and 1,000 randomized repair instances converge); chordal graph coloring [46] assigns disjoint physical qubits to overlapping live intervals (compressing Boyar–Peralta SSA width from 178 virtual wires with 135 wire reuses). (3.2) Under U1–U4, all scratch wires terminate in $|0\rangle$ on $\mathcal{D}$ and net M_X phase polynomials vanish identically in $\mathbb{F}_2[m, x]/I(\mathcal{D})$ (verified across 4/4 multi-round cipher bug injections in 3, 057.9 ms on 10-round AES-128, 4, 249.6 ms on 14-round AES-256, and 9, 315.4 ms on 18-round Keccak-$f[200]$, matching NIST MixColumns `8e4da1bc`); thus the reduced output density matrix is a rank-1 pure state $\mathrm{Tr}(\rho^2_{\mathrm{out}}) = 1$. (3.3, 5.1, 5.2) Verified across 12, 096 semiring triples, sequential grade addition and parallel max/union preserve $\preceq$, enabling optimal DAG Bellman–Knuth extraction; under 90% heralded erasure, minimum uncorrectable weight satisfies $2t_P + t_E \geq d$. (5.3, 5.4, 5.5) Solving $V_1 c(p/p_{\mathrm{th}})^{(d_1+1)/2} \leq \delta$ against $V_0 c(p/p_{\mathrm{th}})^{(d_0+1)/2} \leq \delta$ yields $d_0 - d_1 \geq 2k$; for group-algebra BB-qLDPC matrices over $\mathbb{F}_2[x, y]/(x^\ell - 1, y^m - 1)$, $T_{a,b}$ multiplies both blocks by $x^a y^b$, commuting with polynomial check matrices across all 144 shifts of [[288, 12, 18]]. (5.6, 3.4) In Early-MBU DQI on $H_2(\mathbb{F}_4)$, for any single-symbol error $e_j \in \mathbb{F}_4^*$ at rational point $P_j = (x_j, y_j)$, the syndrome satisfies $(s_0, s_1, s_2) = (e_j, x_j e_j, y_j e_j)$, so $(s_1 \oplus x_j s_0 = 0) \wedge (s_2 \oplus y_j s_0 = 0)$ fires exclusively at index j and `cond-cz` cancels $(-1)^{\mathrm{Tr}(m_j e_j)}$ across all 25 coset leaders (matching 65.22% BCH(15, 7, 5) and 82.98% OPI width reductions, with 0.7338 Heron r2 routed hardware fidelity). In QBlock, when $U^2 = I$, alternating reflections on $|0\rangle^{\otimes m_{\mathrm{anc}}}$ act within invariant 2×2 Jordan subspaces [42] using a single signal ancilla ($m_{\mathrm{anc}} + 1$, 13 anc. on FeMoCo $d = 16$, 91 `TempANDs` = 22.48× $\mathrm{PREP}^\dagger$ reduction, OpenQASM 3.0 T_2 max error 9.98e-16). □

C EXTENDED HARDWARE, DECODER PACING, AND THREE-MODE TABLES

C.1 External Sub-Microsecond Decoder Pacing Feasibility

In the main body (§3.3), Grade Φ_1 abstracts the QEC control plane solely through a target-profile latency parameter τ_{dec}, keeping LakefillQ independent of custom classical hardware. To confirm empirically that $\tau_{\mathrm{dec}} \leq \tau_{\mathrm{cycle}} = 1,000$ ns is achievable under bursty MBU traffic without altering quantum processor hardware, we evaluated an external 500 MHz multiplierless Braun-tree + Shiloach–Vishkin Union-Find syndrome decoder (7,700 LUT6, 0 DSP48, $F_{\max} = 614.25$ MHz, 340 ns cryo-SerDes / 420 ns Heron r2 framing, 660 ns cycle slack, 79.55% sparse fast-path bypass). Across 80,000 Stim QEC cycles, unpaced software MWPM ($\tau_{\mathrm{dec}} = 4,500$ ns) accumulates a 146,205.0 ns FIFO backlog and drops $d = 5$ Heron r2 logical survival from 0.9973 down to 0.5700, whereas any external decoder satisfying $\tau_{\mathrm{dec}} \leq \tau_{\mathrm{cycle}}$ maintains 0.0 ns backlog.

C.2 Extended Cryptographic, Post-Quantum, and Three-Mode Benchmarks

Across public-key factoring and discrete logarithms [24, 32, 40, 41, 44, 64, 95], LakefillQ spans:

- **RSA-2048 and RSA-4096:** Single-focus coset zipper (1,924 Q), Braun-tree RNS-CRT (896 Q RSA-2048; 1,728 Q RSA-4096), and split-register skew-binary carry-save (1,364 Q RSA-2048;

2,644 Q RSA-4096). Decoupling limb carries preserves surface-code distance $d = 27$ on RSA-2048 (891,348 yoked SC physical qubits in 3.19 superconducting days or 24.9 transversal neutral-atom days, and 344,304 BB-qLDPC physical qubits, vs. 1,755,670 yoked SC qubits under $d = 29$ ripple-carry inflation), while RSA-4096 maps to 1,843,882 yoked SC or 583,248 BB-qLDPC physical qubits.

- **ECC-256:** Packed Bézout + Jump-16 pebbles (1,088 Q, 288,216 BB-qLDPC physical qubits, 11,657,520 signed odd-window `TempANDs`; 17,280 single-module Q102 ions vs. 10,490 ions under $K = 4$ serial multishot folding).
- **Symmetric, Hash, Lattice, and Code-Based Workloads:** Pareto in-place Boyar–Peralta AES-128/192/256 (290 / 354 / 418 Q, 392 Q pure in-place, 17,000 AES-256 `TempANDs`), `DLog-5` min-width (261 Q / 325 Q / 389 Q, 26,264 AES-128 Grover Toffolis), bidirectional half-cipher MBU Grover (6,527 AES-128 `TempANDs`, 7,392 Block-2 cone `TempANDs`), SLH-DSA-128s/256s (4,320 / 4,064 Q with 41.56% constant-lane sponge folding), Classic McEliece-348864 (7,744 Q, 814,336 forward Toffolis, 5.24× QM4RI speedup), BIKE-L1 (36,969 Q, 109.24× DOOM rotational amortization), ML-KEM-768 (3,072 Q, 322,560 Toffolis), ML-DSA-65 (5,888 Q, 1,224,704 Toffolis), and OPI DQI (104,448 → 9,984 Q).

Table 10 summarizes LakefillQ's three-mode execution results across Single Machine, Parallel Machine ($M_{\max} = 256$ under Rule T-ParGuard, verified on `distributed_qrns_shor_u4.lfq` with zero reverse ebits and noisy optical DEJMPS distillation at $d = 3, p = 0.05$ consuming 2.11M raw pairs across 131 steps), and Serial Multishot regimes (including noisy Heron r2 Hubbard QPE: 0.7540 Ha Single Machine error vs. 0.0053 Ha 4-QPU Parallel Machine and 0.0827 Ha sampled / 0.0795 Ha exact Serial Multishot at 6.10× per-shot depth reduction).

Table 10. Three-mode compilation summary across Single Machine, Parallel Machine ($M^* \in \{1..256\}$, $M_{\max} = 256$), and Serial Multishot execution.

Workload	Single Q	Per-Run `TempANDs`	Optimal $M^* \times Q_{\text{loc}}$	Optical Ebits	Multishot / Notes
AES-128 Grover	261	6,527	1×261 (Zalka)	0	Bidirectional MBU
ECC-256 ECDLP	1,088	614,040	4×312	9,744 logical	10,490 Q102 ions ($K = 4$)
RSA-2048 Shor	1,364	36,875,000	6×268 (QRNS)	934,200 raw	0-carry QRNS + U4 GHZ
RSA-4096 Shor	2,644	142,812,500	8×356 (24×228)	0 reverse ebits	Knee at $Q_{\text{cap}} = 230$
FeMoCo THC QSVT	2,142	85,156,250	$64 \times 2,142$ (shots)	0 (shot-par.)	6,890 `TempAND`/walk

D OPEN-SOURCE BUG CATALOG, MUTATION MATRIX, AND LEAN 4.35.0 INDEX

Open-Source Audit and Seven Projected Guards. Seven general static guards projected from $\Gamma_{\text{in}} \vdash t : \boldsymbol{Q}_e(A) \dashv \Gamma_{\text{out}}$ (Tables 11 and 12) catch all 14 latent correctness bugs across the nine audited repositories, including 20/32 corrupted basis rows from assuming $\chi_5^2 = I$ in ReVerC, missing PLU permutation uncomputation (4.23 ms XAG check), 7.36× QuantumISD skip-rightmost undercounting (94.53% guarded Aer success probability), and 1.1796 spectral error from applying the degree-2 Chebyshev walk shortcut when $U^2 \neq I$ in Qualtran THC. Full diagnostic and Redex fuzzer artifacts are recorded in `lakefillq_holistic_language_cohesion_report.json` and `lakefillq_redex_surface_diagnostic_fuzzer_report.json`.

Lean 4.35.0 Mechanization Index (2,305 Theorems Across 82 Modules). Every formal claim in this paper maps to zero-`sorry`, zero-`axiom` Lean 4.35.0 theorems verified by the Lean kernel:

- **Core Calculus & Duoidal Capabilities (§3, Thms. 3.1–3.3, 3.4):** `lakefill_qsub_qphys_architecture_master_closure_theorem`, `asymmetric_within_apply_mbu_`

Table 11. Detailed breakdown of representative bugs uncovered across nine open-source quantum repositories.

Repository Family	Bug Pattern	Empirical Failure Witness	Static Guard
ECC / Kaliski / Factoring	$x = 0$ inversion; $\mathbb{F}_{13}$ offset; coset carry	Dirty ancilla / wrap-around	T-DomImg, U2
ReVerC / T-Type AES	$\chi_5^2 \neq I$ (20/32 rows); missing PLU perm	Corrupted state (4.23 ms catch)	T-DomImg, T-BorrowZ
QuantumISD / DQI	Singular QGJE (7.36×); wt-> t syndrome	Amplitude leak (94.53% fixed)	T-DomImg, T-CounitU4
Qualtran / Cobble	$d = 2$ walk on $U^2 \neq I$; PREP† sign	1.1796 op. error; LCU cancel	T-SubNorm, T-PhaseClean

Table 12. Correspondence between the seven coordinate projections of $\Gamma_{\text{in}} \vdash t : Q_e(A) \dashv \Gamma_{\text{out}}$ and the standalone Python OpenQASM 3.0 / Qiskit linter (`ae/lakefillq_lint.py`, 16/16 recall, 0% false positives).

Judgment Projection	Static Obligation Enforced	Wild Bugs Prevented
T-DomImg (U1)	Domain-restricted inverse $g(x, f(x)) = x$ on $\mathcal{D}$	Kaliski $x = 0$, ReVerC χ_5, singular QGJE, DQI
T-BorrowZ	Read-only BorrowZ context non-interference	In-place PLU alias, cold-limb overwrite
T-PhaseClean (U2)	Canonical discharge & $\Phi_X \equiv 0$ before H/QFT	Uncollapsed carry rail, PREP† sign flip
T-CounitU4 (U4)	Bilinear $\mathbb{F}_2$ M_X phase cancellation at exit	Omitted cond-cz, remote GHZ phase leak
T-SubNorm (QBlock)	$U^2 = I$ involutory guard & α, ε bounds	Qualtran $d = 2$ Chebyshev walk corruption
T-QubitCap (Φ_1–Φ_3)	Hardware capacity, τ_{dec} pacing, exclusion	Lindley backlog stall, blockade crosstalk, hook
T-ParGuard ($r \triangleright e$)	Convex M^* selection with Zalka $\sqrt{M}$ penalty	Suboptimal Grover/ISD multi-QPU splitting

Table 13. Compiler soundness mutation testing: all 12/12 (100.0%) injected compiler bugs are killed by simulator counterexamples.

Mutant Group	Disabled Compiler Check	Triggering Counterexample	Killed
M1–M4 (QSub)	Disable U1 domain, U2 canonical, U3 branch, U4 taint	$\text{Tr}(\rho^2) \leq 0.50$ / norm leak	4/4
M5–M6 (SSA/XAG)	Ignore chordal live overlap; skip XAG cut-sweep	Recycled-wire clobber; 10-rd taint	2/2
M7–M9 (QPhysGrade)	Ignore Φ_1 τ_{dec}, Φ_2 blockade, Φ_3 hook flag	Backlog divergence; $d_{\text{eff}} : 5 \to 3$	3/3
M10–M12 (QBlock/M^*)	Skip $U^2 = I$; drop Zalka $\sqrt{M}$; drop reverse GHZ U4	1.18 spec. err.; Bell-cut blowup	3/3

`halves_t_count_zero_phase_taint`, `thm22b_f28_chordal_register_packing_clean_ancilla_safety`, `elaborateToKern_zero_residual_forms`, `openqasm3_mbu_feedforward_sound`, `repairPhysGrade_restores_wellTyped`, `rwire_ordered_duoidal_interchange`, `rwire_borrow_then_move_safe_vs_par_conflict`, `deltaZ_ccx_and_tempAnd_contraction_sound`, `sequential_callQ_chordal_scratch_saves_wires`, `auto_cluster_profile_eliminates_cost_drift`, `qsub_logical_preservation_and_unit_purity_soundness`, `wire_level_symbolic_type_inference_module_closure`, `bilinear_xag_mbu_collapses_arbitrary_degree_to_zero`, `multiround_mbu_session_total_quantum_phase_unit`, `cut_sweeping_local_unsat_implies_global_multiround_soundness`, `anf_to_xag_escalation_preserves_zero_verdict`, `qphys_grade_seq_associative_and_identity`, `qphys_phi1_phi2_phi3_strict_independence_witnesses`, `lcu_projection_sound_under_u1_u4`, `qsvt_chebyshev_ancilla_compression`, `berryGidney_femoco_L2048_91_lt_2046`, `linTong_serial_multishot_ancilla_and_phi1_advantage`.

- **Uncomputation & Quantum Data Structures (§4, Props. 4.1, 4.2, 4.3–4.11):** `degree_guarded_hmr_cz_soundness_and_cubic_rejection`, `catch1_gidney_mbu_cannot_zero_out_macro_pebble_windows`, `reversible_tree_pebbling_simultaneous_2h_plus_1_vs_sequential_h_plus_1`, `api1_huet_focus_zipper_contract_and_cost`, `api2_skew_binary_carry_save_contract_and_cost`, `skew_binary_canonical_`

discharge_preserves_exact_integer_value, api3_braun_tree_fold_contract_and_cost, braun_tree_size_and_lsb_routing_invariant, api4_packed_bezout_shift_zipper_contract_and_cost, hua_identity_with_singular_zero_one_correction_exact, bezout_half_domain_folding_strict_cofactor_bijection, api5_binomial_pebble_schedule_contract_and_cost, api6_monoidal_finger_tree_simhash_contract_and_cost, api7_bruhat_pivoted_lu_and_streaming_syndrome_contract_and_cost, bruhat_lu_f2_zero_pivot_requires_permutation_and_recovers_all_gl4, api8_canonical_critbit_bitonic_dictionary_contract_and_cost, parallel_walker_bitonic_74_47x_amortization_and_width_tradeoff, api9_walker_vose_alias_table_contract_and_cost.

- **ZX e-Graph, Hardware Step-Down, & AG DQI (§5–§6, Lemmas 5.1–5.2, Prop. 5.5, Thms. 5.3–5.6, Cor. 5.4):** zx_spider_fusion_mod8_sound, zx_t_tdg_gadget_cancel, zx_t_t_fuses_to_clifford_s, zx_cx01_parity_transport_sound, shared_twin_ccz_zx_8t_eq_ideal, hermitian_h2_all_points_on_curve, hermitian_h2_exceeds_reed_solomon_bound, hermitian_h2_weight1_syndrome_injective, hermitian_locator_recovers_exact_coset_leader, hermitian_early_mbu_law_u4_phase_cancel, lindley_zero_backlog_when_paced, lindley_linear_backlog_when_unpaced, external_decoder_pacing_concrete_separation, rccx_paired_within_and_phase_cancellation, knuth_dijkstra_optimal_acyclic_egraph_extraction, four_target_pareto_separation_and_qroam_dominance, bernoulli_scaled_survival_verified, mixed_pauli_erasure_min_fault_weight, erasure_doubles_correctable_faults_all_odd_distances, volume_reduction_steps_down_code_distance_by_2, rsa2048_and_ecc256_code_distance_step_down_verified, f41_s21_cat_qubit_z_bias_preserving_and_76x_compression, dejmps_distilled_epr_photonic_deque_zipper_threshold_soundness, qldpc_bivariate_bicycle_automorphism_zipper_vs_bridge_catch, berry_gidney_sublinear_sqrt_n_temp_and_dominance, kawamura_law_u2_quotient_erasure_and_canonical_representation, early_mbu_phase_cancellation_algebraic_identity, spanning_tree_peel_commutation_Hc_eq_s, lakefillq_grade_phi1_lindley_zero_backlog, unpaced_mwpm_lindley_backlog_divergence.
- **Diagnostics, Fuzzing, Audits, & Distributed M^* (§6):** complete_12_of_12_compiler_mutant_kill_rate_soundness, multipass_pipeline_preserves_caller_span, buildProvenanceChain_preserves_user_anchoring, chordal_resugar_recycled_step_blames_second_owner, resugar_meas_eliminates_all_compiler_gensyms, minimal_monomial_witness_evaluates_true, undeclared_wire_rejected_outside_watermark, callQ_aliased_mutated_port_rejected, lisp1_dual_position_operator_preserves_gate, surface_adjoint_involutive, braun_seq_invert_exact_identity, enriched_diag_total_3span_hygiene, ecc_f13_offset_trajectory_sound_vs_reverse_uncompute_hazard, kaliski_zero_input_unguarded_dirty_vs_guarded_clean, coset_carry_runway_non_wraparound_and_tightness, reverc_chi5_not_involution, plu_missing_perm_unsound, sphincs_chi_constant_lane_eliminates_two_ands_per_row, qgje_non_singular_guard_soundness, mceliece_348864_qm4ri_tempand_reduction, bike_l1_doom_amortized_109x_speedup, qwalk_d2_equals_chebT2_iff_involution, wx_to_refl_phase_d2_sign_flip_cancels_lcu_sum, femoco_thc_law_u4_walk_cost_strictly_dominates_baselines, t_par_guard_zalka_fallback_soundness, t_par_guard_bell_cut_overhead_triggers_single_machine_fallback, law_u4_ghz_cat_disentangle_halves_lookup_ebits, qrns_distributed_zero_carry_ebits_soundness, qpar_m_

star_ucurve_knee_pruning_and_monolithic_fit, qpar_m256_horizon_rsa4096_knee24_and_optical_tree_scaling, gf16_inv_4and_obstruction_and_5and_exactness, bidirectional_half_cipher_mbu_grover_dominance, mitm_leaf_gadgets_shor_and_isd_dominance, dlog5_inplace_aes_sbox_256_exact_and_5ancilla_cleanup, ecc256_dual_qec_and_m16_cluster_zalka_contrast, t_par_guard_grover_fallback_soundness, distributed_law_u4_lookup_qrns_and_m_star_soundness.